

RAGalyst: Automated Human-Aligned Agentic Evaluation for Domain-Specific RAG

Joshua Gao*, Quoc Huy Pham*, Subin Varghese, Silwal Saurav, Vedhus Hoskere

*Equal Contribution

University of Houston

<https://joshuakgao.github.io/RAGalyst>

Abstract—Retrieval-Augmented Generation (RAG) is a critical technique for grounding Large Language Models (LLMs) in factual evidence, yet evaluating RAG systems in specialized, safety-critical domains remains a significant challenge. Existing evaluation frameworks often rely on heuristic-based metrics that fail to capture domain-specific nuances and other works utilize LLM-as-a-Judge approaches that lack validated alignment with human judgment. This paper introduces RAGalyst, an automated, human-aligned agentic framework designed for the rigorous evaluation of domain-specific RAG systems. RAGalyst features an agentic pipeline that generates high-quality, synthetic question-answering (QA) datasets from source documents, incorporating an agentic filtering step to ensure data fidelity. The framework refines two key LLM-as-a-Judge metrics—Answer Correctness and Answerability—using prompt optimization to achieve a strong correlation with human annotations. Applying this framework to evaluate various RAG components across three distinct domains (military operations, cybersecurity, and bridge engineering), we find that performance is highly context-dependent. No single embedding model, LLM, or hyperparameter configuration proves universally optimal. Additionally, we provide an analysis on the most common low Answer Correctness reasons in RAG. These findings highlight the necessity of a systematic evaluation framework like RAGalyst, which empowers practitioners to uncover domain-specific trade-offs and make informed design choices for building reliable and effective RAG systems. RAGalyst is available on our Github.

Index Terms—Domain-Specific, Retrieval-Augmented Generation, LLM-as-a-Judge, Synthetic Dataset Generation, Question-Answering Dataset, RAG Evaluation.

I. INTRODUCTION

Although modern Large Language Models (LLMs) are great synthesizers of information, they still suffer from hallucinations [1], [2], which refers to the generation of content that appears plausible but is factually incorrect or unsupported by evidence. Mitigating hallucinations is especially important in safety-critical applications (e.g., military operations, cybersecurity, and bridge engineering) where inaccurate information can lead to serious consequences and undermine trust in artificial intelligence (AI) systems [3], [4]. Retrieval-Augmented Generation (RAG) has been widely adopted to mitigate hallucinations by grounding responses in provided context [5], [6].

A key advantage of RAG is its ability to provide models with dynamic, inference-time access to private and domain-relevant documents [5], [7]. However, leveraging this ability is highly sensitive to several domain-specific variables.

Source documents in specialized fields often include out-of-distribution content—such as dense jargon or unconventional formatting—that lies outside the training corpora of the LLM and embedding model, hindering their ability to interpret the retrieved text. Similarly, the optimal text chunking strategy and retrieved context lengths may differ in different domains due to document structure and the amount of contextual information required to synthesize information. For example, bridge engineering documents may require an understanding of deterioration trends across extended spans of inspection narratives and historical measurements, favoring longer context windows to capture cross-report dependencies. On the other hand, cybersecurity documents tend to present short but information-dense logs (e.g., a few lines of packet capture) may be sufficient for accurate interpretation, making smaller chunks preferable.

The challenge of adapting RAG systems to specialized domains is compounded by the difficulty of accurately evaluating their performance. Such evaluation requires domain-specific benchmarks and metrics that produce reliable results.

Given the widespread need for RAG systems and the innumerable domains in which they may be used, manually producing benchmark datasets for tuning of parameters can be very challenging. State-of-the-art RAG evaluation frameworks increasingly employ LLMs to generate synthetic Questions and Answers (QAs) from a knowledge base that are then used to evaluate RAG system performance. The quality of these datasets then become critical in the reliability of reported evaluations. Most approaches require human validation in combination with heuristics to ensure dataset quality. Other fully automated generation pipelines like RAGAS [8] lack rigorous quality filtering, resulting in an unreliable QA dataset.

Beyond dataset generation, evaluation metrics are also vital to directly estimate RAG system performance. Early works in RAG evaluation that often rely on traditional heuristic metrics, such as Bilingual Evaluation understudy (BLEU) [9] and Recall-Oriented Understudy for Gisting Evaluation (ROUGE) [10], operate by measuring the literal overlap of words or phrases between the generated answer and a reference answer. A challenge lies in their inability to capture semantic meaning; a factually correct response that uses different phrasing would be unfairly penalized, while a nonsensical answer that shares keywords might score deceptively high. To overcome

the rigidity of these lexical metrics, researchers have begun leveraging LLM-as-a-Judge to better evaluate the semantic nuances of generated answers [8], [11]–[15]. However, these LLM-based metrics have not yet been rigorously examined for human-alignment. There is a need to develop human-alignment mechanism to LLM as a judge metrics that can help more reliably automate benchmarking and tuning of domain-specific RAG systems.

This paper introduces RAGalyst, an end-to-end human-aligned agentic framework for domain-specific RAG evaluation. Our framework integrates refined LLM-as-a-Judge metrics and fully automated agentic QA benchmark construction to enable rigorous benchmarking and more reliable deployment of RAG systems across diverse domains. We evaluate the human-alignment procedure on two metrics, the Answerability metric - used in synthetic QA generation, and the Answer Correctness metric - used in RAG answer evaluation.

The main contributions of this paper are as follows:

- We present an automated agentic framework for evaluating domain-specific RAG systems, integrating document preprocessing, reliable synthetic QA generation, embedding model and LLM-based evaluation modules.
- We introduce novel human benchmark-aligned prompt-optimized LLM-as-a-Judge formulations for RAG metrics, namely the Answerability and Answer Correctness metrics, which outperform state-of-the-art formulations as judges.
- We evaluate our framework and demonstrate applicability on documents from three different specialized domains, namely military operations, cybersecurity, and bridge engineering.

II. RELATED WORK

A. Retrieval-Augmented Generation (RAG)

RAG addresses hallucinations by integrating a retrieval module that dynamically retrieves relevant textual chunks from a knowledge base during inference [5]. Early RAG systems followed a straightforward retrieve-then-generate pipeline [16]. Recent advances include modularization for flexible component composition [17], adaptive retrieval strategies that dynamically adjust retrieval depth [18], and graph-based reasoning mechanisms for multi-hop information synthesis [19]. The Atlas model demonstrated that few-shot learning achieved strong performance even with relatively small parameter counts when using a dense retriever and joint pre-training strategies [7]. LongRAG uses large context chunks and leverages long context LLMs to reduce retrieval noise and improve semantic integrity [20]. On the other hand, OP-RAG argues that naive use of long-context models may dilute relevant content, and that order-preserving RAG techniques can offer superior efficiency and answer quality by maintaining the original document structure during chunk selection [21].

B. Domain-Specific RAG

While general-purpose RAG systems demonstrate promising capabilities in knowledge-intensive tasks, they often under-

perform in domains where specialized knowledge or jargon is involved. Domain-specific RAG bridges this gap by tailoring both retrieval and generation processes to the target domain. RAFT [22] proposes Retrieval-Augmented Fine-Tuning, a training procedure that fine-tunes LLMs to handle bad retrieval to enhance robustness and improve downstream QA performance on specialized datasets such as PubMed and HotpotQA. Similarly, Li et al. [6] created a synthetic dataset sourced from Carnegie Mellon University’s website, and proposed a domain-specific RAG pipeline. Nguyen et al. [23] show that fine-tuning both the embedding model and the generator significantly improves performance on complex datasets like FinanceBench.

To benchmark and evaluate RAG capabilities in expert domains, Wang et al. [24] introduced *DomainRAG*, a comprehensive benchmark for Chinese university enrollment data. Their work identified six critical abilities for domain-specific RAG systems: conversational handling, structural understanding, denoising, multi-document reasoning, time sensitivity, and faithfulness to external knowledge. They demonstrated that current LLMs struggle significantly in a closed-book setting, validating the necessity for domain-specific RAG systems.

C. RAG Evaluation

RAG performance depends on a complex interplay of components, including retrieval module, document chunking strategies and model prompting [25]. Unlike standalone LLMs, RAG systems are highly sensitive to changes in these components, making generalization between tasks and settings extremely difficult.

Traditional evaluation metrics such as BLEU [9], ROUGE [10], or exact match fail to account for the modular and domain-specific nature of RAG pipelines. Consequently, a variety of frameworks have emerged to address this challenge. RAGEval [11] and ARES [12] are prominent reference-free and semi-supervised evaluators that assess context relevance, answer faithfulness, and informativeness without requiring ground-truth answers. ARES, in particular, offers statistical confidence bounds and modular component scoring, enabling accurate diagnostics even across domains.

End-to-end evaluation often obscures specific failure points within the RAG pipeline (such as suboptimal chunking, inaccurate retrieval, ineffective reranking, or hallucinated generation) making it difficult to isolate which component is responsible for performance degradation. To address this, eRAG [26] proposes a document-level relevance scoring based on LLM output, which correlates better with downstream QA performance than traditional query-document metrics. CoFE-RAG [27] expands this perspective by evaluating all stages: chunking, retrieval, reranking, and generation. These techniques allow developers to diagnose failure sources within the RAG pipeline and improve system interpretability. The RAGAS framework [8] offers a comprehensive and modern evaluation pipeline, encompassing dataset generation from documents, LLM-based metrics, and a modular evaluation architecture. As it is the only actively maintained framework of its kind, we adopt RAGAS as the primary baseline for

evaluating the performance of our proposed method. Even though all of these end-to-end evaluation frameworks have made significant strides, they rely on some degree of manual validation for QA dataset generation. Moreover, their metrics have notable limitations: rule-based metrics often fail to capture subtle semantic nuances, while works that use LLM-based metrics are rarely benchmarked for alignment with human judgment. As a result, such metrics may not fully reflect the intended evaluation objectives or agree with human assessments.

D. Tuning LLM-as-a-Judge Evaluation

Recent advances in LLM evaluation have increasingly focused on methods that improve alignment with human judgment and reduce variability in scoring. AutoCalibrate [28] auto-tunes the evaluation prompts for better human alignment, ensuring that the scores reflect the actual preferences of the user. DSPy [29], [30] is a declarative framework that enables the programmatic creation and refinement of prompts for LLMs. Evaluation frameworks such as PoLL [31] advocate a panel of various LLM evaluators to reduce bias and variance in generation scoring. These are powerful methods that have yet to be applied in RAG evaluation.

E. Synthetic QA Data Generation

The evaluation of RAG systems typically relies on a question-answering (QA) datasets. However, such QA datasets are often unavailable or insufficient in specialized domains [32]. This limitation has driven a surge in research focused on generating high-quality synthetic QA data.

Alberti et al. [33] pioneered a round-trip consistency approach that combines answer extraction, question generation, and answer re-verification to create high-confidence Question-Answer-Context (QAC) triplets. Shakeri et al. [34] proposed an end-to-end transformer-based generator that outputs both the question and the answer from a given passage.

Recent QA data generation has increasingly shifted toward LLM-based approaches. Bai et al. [32] tailored prompt engineering and summarization strategies to generate more challenging clinical QAC triplets from EHRs. Bohnet et al. [35] leveraged long-context LLMs to generate QAC triplets over entire books, demonstrating that automatic generation can rival or surpass human-crafted data sets in narrative domains. RAGAS [8] leverages agentic LLM designs and knowledge graph to generate context-rich and diverse question–answer pairs, with the LLMs simulating multiple personas during the generation process.

While these synthetic QA generation frameworks are impressive, none of them employ a fully automatic end-to-end pipeline that leverages LLM-based metrics to assess QA quality. RAGAS attempts to accomplish this yet their QA generation underperforms even on their own native metrics.

III. METHODOLOGY

This section presents an agentic framework for evaluating RAG systems in domain-specific contexts. The framework

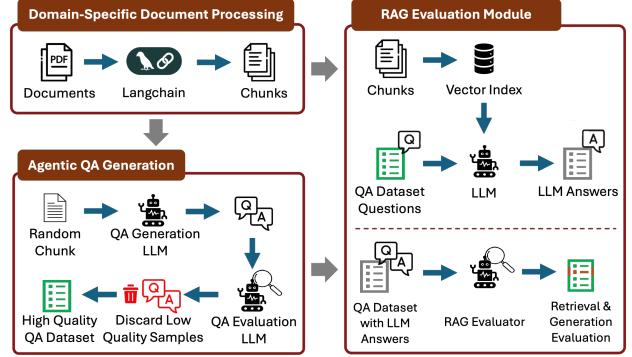


Fig. 1. Overview of the RAGalyst framework that consists of three modules: a pre-processing module that transforms domain-specific documents into text chunks, a QA generation pipeline for producing synthetic question–answer–context datasets, and an evaluation module for assessing RAG system performance.

includes a document preprocessing toolkit, agentic QA generation pipeline, and a set of LLM-based evaluation metrics. Figure 1 illustrates the overall framework.

A. Domain-Specific Documents Preprocessing

The preprocessing of domain-specific documents is the first and most critical step in this RAG framework, as it directly supports both the QA generation and evaluation stages. Without careful handling of document formatting and structure, the quality of retrieval and generated responses suffer significantly.

The framework first leverages LangChain¹ to parse PDF, markdown, and plain text and then divides the documents into smaller units, called chunks, using a token-based chunking strategy. The size of these chunks plays a vital role in downstream performance. Chunks that are too small may lack context, leading to incomplete retrievals. On the other hand, chunks that are too large can dilute relevance and exceed the model context limits. Chunk size significantly impacts retrieval performance across domains due to differences in document structure and content. Optimal chunk size can vary by more than 20% depending on the domain, influencing both precision and recall metrics in RAG evaluations [36]. Since there is no standardized text chunking strategy, we follow OpenAI’s file search tool.² and select a default of a maximum 800 tokens per chunk with an overlap between chunks of 400 tokens. Afterwards, the chunks are vectorized with the selected embedding model and are stored in a vector database.

B. Agentic QA Generation Pipeline

A critical requirement for evaluating RAG systems is the availability of high-quality annotated datasets. These datasets must contain:

- 1) Domain-relevant questions.
- 2) Ground-truth answers.
- 3) Context used to answer each question.

¹<https://www.langchain.com/>

²<https://platform.openai.com/docs/assistants/tools/file-search>

TABLE I
SUMMARY OF EVALUATION METRICS FOR DATASET GENERATION AND RAG EVALUATION

Metric Name	Use	Scale	Used In	Framework
Answer Correctness	Evaluates how accurately the generated answer matches the ground truth	0–1.0	RAG Evaluation	Ours
Answerability	Evaluates whether the question can be answered using only the provided context	0 or 1	Dataset Generation	Ours
Faithfulness	Measures how well the answer is grounded in the provided context	0–1.0	Dataset Generation, RAG Evaluation	RAGAS
Answer Relevance	Assesses whether the answer directly and meaningfully addresses the question	0–1.0	Dataset Generation, RAG Evaluation	RAGAS
Recall@K	Measures whether the ground truth context appears in the top- K retrieved documents	0 or 1	Retrieval Evaluation	Standard
MRR	Computes the inverse rank of the first correct retrieval to evaluate re-ranking quality	0–1.0	Retrieval Evaluation	Standard

However, for many classified or sensitive domains, such annotated datasets are non-existent. Manual annotation of such datasets is often impractical due to confidentiality restrictions, inconsistent formatting, and the high cost of human labeling.

To address this gap, the agentic QA generation pipeline leverages LLMs in a role-driven, autonomous fashion to emulate users to produce context-based, high-fidelity question-answer pairs based on preprocessed document chunks. Each generated answer is grounded directly in its source context and quality is validated using multiple LLM-based metrics from the evaluation module (more details in III-C) to ensure alignment and quality. This approach eliminates the need for labor-intensive annotation while ensuring consistency, reproducibility, and scalability across various domains.

The pipeline operates in three main steps:

a) *Context Sampling*: Chunks generated during the document chunking step are randomly sampled and used as reference contexts for question, answer, context triplet (QAC).

b) *Prompting and QA Generation*: An agentic LLM assumes the role of a user to generate a question that is related to and answerable by the sampled context. The questions is evaluated to ensure it is specific and unambiguous. Then, another agent assumes the role of a subject matter expert to answer the question to generate the ground truth answer.

c) *Validation and Filtering*: To ensure quality, the LLM-based evaluation module validates the generated QA pairs using Answerability, Faithfulness and Answer Relevance metrics (more information in Section III-C2). If the sample is unable to meet the thresholds, it is discarded to preserve dataset integrity. These thresholds are hyperparameters that control quality strictness. The higher values enforce stronger filtering but result in longer generation times due to an increased number of candidate QAs being discarded.

C. RAG Evaluation Module

The RAG evaluation module leverages LLM-as-a-judge to enable automated, scalable, and consistent scoring across

diverse settings. Additionally, the framework also provides essential heuristic metrics for evaluating retrieval performance.

1) *Retrieval Evaluation*: Mean Reciprocal Rank (MRR) and Recall@K are two standard information retrieval metrics used to quantify retrieval effectiveness [25].

MRR computes the average of the reciprocal ranks of the first relevant document across queries, effectively measuring how early the correct context appears in the retrieval list. MRR is particularly well-suited for assessing re-ranking strategies in RAG systems, as it emphasizes placing relevant information as close to the top of the ranked list as possible.

MRR is defined as:

$$\text{MRR} = \frac{1}{|Q|} \sum_{i=1}^{|Q|} \frac{1}{\text{rank}_i}$$

where:

- $|Q|$ is the total number of queries.
- rank_i is the position of the ground truth context in the ranked list for the i -th query.

Recall@K, also referred to as Hit Rate@K, is the average recall across multiple queries where ground truth context appears within the top- k retrieved results for a given query. For each query, recall is a binary metric. It returns 1 if the ground truth context is found within the top k or 0 if not. This metric is particularly important for RAG systems since it calculates the chance the ground truth context is included among the retrieved top candidates.

Recall@ k is defined as:

$$\text{Recall}@K = \frac{1}{|Q|} \sum_{i=1}^{|Q|} \mathbf{1}(\text{rank}_i \leq k)$$

where:

- $|Q|$ is the total number of queries.
- rank_i is the position of the ground truth document for the i -th query.

- $\mathbf{1}(\cdot)$ is the indicator function, which returns 1 if the condition inside is true, and 0 otherwise.
- k is the cutoff rank threshold (for example, Recall@3 considers the top 3 results).

2) *LLM-Based Evaluation*: To ensure the quality of the synthetic QA dataset and support downstream RAG evaluation, the framework combines established metrics with a novel LLM-based evaluation approach. While standard metrics from the RAGAS framework [8] are effective for detecting hallucinations and assessing answer relevance, they fall short in evaluating the quality of generated questions and the correctness of RAG-generated answers. To address these limitations, the framework introduces custom LLM-as-a-judge metrics, which leverage language models to better account for linguistic variation, paraphrasing, and latent knowledge [37]–[40]. This LLM-based strategy enables a more robust, human-aligned evaluation pipeline. The key metrics introduced are:

- **Answer Correctness**: A custom LLM-as-a-judge metric that compares a generated answer to the reference (ground truth) answer. The LLM scores semantic alignment on a continuous scale, offering both a numerical score and a rationale. This enables flexible evaluation even when surface-level phrasing differs.
- **Answerability**: Assesses whether a generated question is fully supported by the retrieved context, without relying on external knowledge. This enforces quality in our synthetic QA dataset generation.

Table I summarizes the definitions of all metrics and their roles within our framework.

IV. FRAMEWORK VALIDATION

This section validates the Agentic Framework for Domain-Specific RAG Evaluation (RAGalyst) by assessing its accuracy, reliability, and practical suitability for domain-specific document tasks. The primary objective is to determine whether the framework’s evaluation metrics and data generation methods are effective enough to support real-world use cases.

First, we assess the reliability of the LLM-as-a-judge evaluation metrics Answer Correctness and Answerability—along with prompt-optimized variants. These metrics are tested against human-annotated references to ensure alignment with human judgment.

Second, the synthetic QA datasets generated by the agentic QA generation pipeline are evaluated against publicly available domain-specific QA sets, and datasets generated by RAGAS, to assess its fidelity and reliability in evaluating RAG systems. This comparison aims to verify whether the synthetic data produced by this framework serves as a viable substitute for manually curated datasets in high-stakes domains.

A. Metrics Validation

We evaluate the performance of both Answer Correctness and Answerability by computing the Spearman correlation between their scores and human annotations. To calculate the

standard error of a Spearman correlation, we use Bonnett and Wright [41] standard error (SE) approximation:

$$SE(\rho_s) = \sqrt{\frac{1 + \frac{\rho_s^2}{2}}{n - 3}}$$

Where:

- $SE(\rho_s)$ is the standard error of the Spearman correlation.
- ρ_s is the expected Spearman correlation coefficient.
- n is the sample size.

Since a manually crafted Answer Correctness and Answerability prompt is unlikely to be optimal, we turn to DSPy’s [29], [30] COPRO, MIPROv2, and LabeledFewShot optimizers. These optimizers systematically refine prompts using supervision and feedback signals, improving alignment with human annotations.

COPRO and MIPROv2 are automatic instruction optimizers where the instructions in the prompt are tweaked and validated. The COPRO optimizer generates and refines new instructions for each step, and optimizes them with coordinate ascent on a defined metric function. MIPROv2 generates instructions and few-shot examples in each step of optimization. The instruction generation is data-aware and demonstration-aware, and uses Bayesian Optimization to effectively search over the space of generation instructions/demonstrations.

LabeledFewShot is a automatic few-shot learning optimizer. It randomly samples k examples from a labeled dataset and adds them to the prompt as demonstrations.

1) *Answer Correctness*: This metric was benchmarked against the Semantic Textual Similarity Benchmark (STS-B), part of the GLUE benchmark suite³. STS-B consists of sentence pairs drawn from sources such as news headlines, image captions, and online forums, with each pair annotated by humans with a similarity score ranging from 0 (completely dissimilar) to 5 (semantically equivalent). To adapt the dataset to the evaluation task, the first sentence in each pair was treated as the ground-truth answer and the second as the model-generated answer. The original STS-B scores were normalized from a range of 0–5 to a range of 0.0–1.0 to align with the output of our LLM-based Answer Correctness metric.

In this validation, we compare our proposed Answer Correctness metric—along with a prompt-optimized variant—against baseline methods including cosine similarity and RAGAS’s Answer Correctness. Both RAGAS and our framework utilize the GPT-4o-mini model at a temperature of 0 to ensure a consistent evaluation backbone across approaches, but we also compare against other state of the art LLMs. Cosine similarity is computed using the Qwen3-Embedding-8B model, the top-ranked model on the MTEB leaderboard⁴ for the STS task at the time of writing.

On a test set of 500 randomly sampled STS-B sentence pairs, our Answer Correctness metric achieves a mean Spearman correlation of 0.874 with a standard error of 0.053 with GPT-4o-mini. This outperforms both Cosine Similarity ($\rho_s =$

³<https://gluebenchmark.com/tasks/>

⁴<https://huggingface.co/spaces/mteb/leaderboard>

TABLE II
HUMAN ALIGNMENT OF METHODS FOR ANSWER CORRECTNESS AND ANSWERABILITY

Method	Model	Answer Correctness ρ_s	Answerability ρ_s
Cosine Similarity	Qwen3-Embedding-8B	0.622	—
RAGAS	gpt-4o-mini	0.836	—
Ours	gemma-3-27b-it	0.862	0.596
Ours	Qwen3-30B-A3B-Instruct-2507	0.851	0.605
Ours	gemini-2.5-flash-lite	0.849	0.665
Ours	gemini-2.5-flash	0.777	0.749
Ours	gemini-2.5-pro	0.805	0.752
Ours	gpt-4o-mini	0.874	0.700
Ours	gpt-4.1-nano	0.857	0.436
Ours	gpt-4.1-mini	0.873	0.670
Ours	gpt-4.1	0.857	0.723
Ours: COPRO Optimized	gpt-4o-mini	0.881	0.670
Ours: MIPROv2 Optimized	gpt-4o-mini	0.887	0.669
Ours: MIPROv2 & LabeledFewShot Optimized	gpt-4o-mini	0.894	—
Ours: LabeledFewShot Optimized	gpt-4o-mini	—	0.632

0.622, SE = 0.049) and RAGAS ($\rho_s = 0.843$, SE = 0.052), demonstrating stronger alignment with human-annotated similarity scores.

We first refine the prompt instructions using both the COPRO and MIPROv2 optimizers on a training set of 500 STS-B samples, and evaluate performance on the test set. We adopt DSPy’s default parameters for both optimizers and use GPT-4o-mini to generate optimized prompt candidates. The metric achieves $\rho_s = 0.881$ (SE = 0.053) with COPRO and $\rho_s = 0.887$ (SE = 0.053) with MIPROv2.

Since the MIPROv2 optimized Answer Correctness metric performs the best, we further optimize its prompt with LabeledFewShot optimizer which pushes performance to $\rho_s = 0.894$ (SE = 0.053) at a k of 8. We ablate k in Figure 2. This combination of the MIPROv2 and LabeledFewShot optimizers performs the best for the Answer Correctness metric, and therefore use this version of the metric throughout the remainder of this paper.

2) *Answerability*: The Answerability metric was validated using the Stanford Question Answering Dataset 2.0 (SQuAD 2.0)⁵. SQuAD 2.0 extends the original SQuAD dataset by including over 50,000 unanswerable questions written adversarially to resemble answerable ones. Each entry in the dataset contains a question, a context passage, a ground-truth answer, and a binary flag indicating whether the question is answerable given the context. This makes it an ideal benchmark for evaluating the Answerability of QACs, as it directly tests whether the provided context alone is sufficient to support a meaningful response. Each LLM evaluates the Answerability of the QACs with a temperature of 0.

Using Gemini 2.5 Pro on a test set of 500 randomly sampled SQuAD 2.0 QA pairs, our Answerability metric achieves a mean Spearman correlation of 0.752 with a standard error of 0.051

We first refine the prompt instructions using both the COPRO and MIPROv2 optimizers on a training set of 500

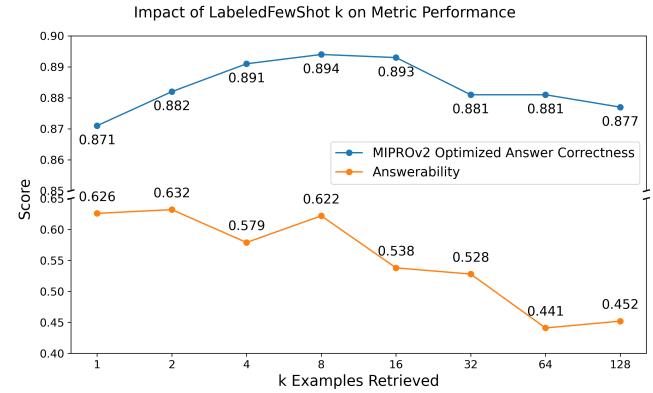


Fig. 2. We ablate both the MIPROv2-optimized Answer Correctness metric and the non-optimized Answerability metric using the LabeledFewShot optimizer. Our results show that Answer Correctness achieves its best performance with 8 examples, whereas LabeledFewShot optimization provides no improvement over our handcrafted prompt for Answerability.

SQuAD 2.0 QA pairs, and evaluate performance on a the test set. We adopt DSPy’s default parameters for both optimizers and use GPT-4o-mini to generate optimized prompt candidates. The metric achieves $\rho_s = 0.670$ (SE = 0.050) with COPRO and $\rho_s = 0.669$ (SE = 0.050) with MIPROv2.

Since neither automatic instruction optimizers improve performance, we apply LabeledFewShot optimizer on the non-optimized Answerability metric with GPT-4o-mini. This achieves $\rho_s = 0.632$ (SE = 0.049) at a a k of 2. We ablate k in Figure 2. All optimizers do not improve performance for the Answerability metric, and therefore use the non-prompt optimized version of Answerability throughout the remainder of this paper.

3) *Metric’s LLM Selection*: The results in Table II suggest that GPT-4o-mini provides the best balance of performance metrics, inference speed, and cost. For this reason, **GPT-4o-mini is used for all LLM-based metrics throughout the remainder of this paper**.

⁵<https://rajpurkar.github.io/SQuAD-explorer/>

B. Agentic QA Dataset Generation Pipeline Validation

To rigorously validate the performance of our Agentic QA generation pipeline, we constructed three domain-specific datasets using our framework. The domains military operations (army), cybersecurity, and bridge engineering (engineering) were selected for their real-world importance and the abundance of diverse, unstructured source material such as technical documents, reference books, and cheatsheets. For each domain, we generated 500 QA pairs, resulting in a total of 1,500 samples that reflect a broad spectrum of question types and contextual difficulty levels.

To establish a fair and consistent baseline for comparison, we used the same source materials to generate three additional datasets using the RAGAS framework. Both frameworks were configured to produce Single Hop-specific QA datasets, which require no or minimal reasoning. This controlled setup ensures that the only variable under evaluation is the generation method itself. We evaluated both approaches using a mixture of our proposed metrics (Answerability) and the RAGAS metrics (Faithfulness, Answer Relevance), allowing us to assess our performance against a strong existing baseline.

In addition to synthetic comparisons, we incorporated a benchmark evaluation against human-annotated datasets. These datasets were selected based on their widespread use in the QA and RAG research communities, their utility in real-world applications, and their sample sizes being large enough to ensure statistical significance.

- The COVID-QA dataset from deepset comprises of 2,019 expert-annotated QA pairs drawn from 147 COVID-19 scientific articles by 15 biomedical specialists [42]. It is highly relevant to real-world biomedical information retrieval tasks.
- The RepLiQA dataset by ServiceNow Research consists of approximately 90,000 question-answer-context triplets, all carefully written and annotated by human experts using fictional yet coherent narratives [43]. It is specifically designed to rigorously evaluate a model’s ability to reason over unseen, non-factual content.

By applying the same evaluation framework to both the machine-generated and human-annotated datasets, we are able to observe how closely our pipeline aligns with human-created content across different domains. This evaluation approach supports a more balanced assessment of the method’s reliability and generalizability.

Table III demonstrates that our framework consistently outperforms RAGAS across both its native metrics and our evaluation criteria. This performance gap stems primarily from two critical limitations in the RAGAS pipeline. First, RAGAS lacks an effective filtering mechanism to exclude low-quality samples, leading to noisier datasets. Second, its use of a single multitask prompt for generating both questions and answers reduces generation quality, aligning with prior findings that multitasking degrades large language model performance [44]. Notably, while the RAGAS-generated dataset achieves reasonable scores in Faithfulness and Answer Relevance, the

dataset’s Answerability remains significantly lower. This pattern suggests that the model generates answers that appear contextually appropriate and responsive but are not explicitly supported by the retrieved evidence, indicating a tendency toward informed hallucination rather than faithful grounding.

Although our custom QA generation pipeline produces higher quality samples, it operates at a slower rate. It generates 100 QA samples at an average of 0.881 samples per minute (single-threaded) and 7.039 spm (16-threaded). In contrast, RAGAS achieves 14.627 spm (single-threaded) and 33.842 spm (16-threaded). This slower generation rate is primarily due to our more stringent QA filtering.

We provide an example generated QA from the Military Operations domain.

Example Generated QA	
Text Chunk as Context:	
Chapter 7 7-6 TC 3-21.76 26 April 2017	
7-29. R&S teams move using a technique such as the cloverleaf method to move to successive OPs. (See figure 7-1.) In this method, R&S teams avoid paralleling the objective site, maintain extreme stealth, do not cross the LOA, and maximize the use of available cover and concealment.	
7-30. During the conduct of the reconnaissance, each R&S team returns to the RP when any of the following occurs:	
– All their PIR is gathered. – The LOA is reached. – The allocated time to conduct the reconnaissance has elapsed. – Enemy contact is made.	
<i>LEGEND: ORP – objective rally point; RP – release point; S&O – surveillance and observation.</i>	
Question: What action must an R&S team take if they make enemy contact during reconnaissance?	
Answer: If an R&S team makes enemy contact during reconnaissance, they must return to the RP.	

V. EXPERIMENTS

To demonstrate the applicability of RAGalyst, this section showcases the experimental evaluation of various RAG approaches using the framework. These experiments assess key components of RAG systems, including the embedding retrieval, LLM generation, and context length. In addition, we analyze the potential bias in LLM’s preference for self-generated text, and we provide an analysis on reasons for low Answer Correctness.

A. Embedding Retrieval Evaluation Across Domains

The performance of any RAG system is fundamentally dependent on the quality of its embedding model. The embedding model determines how well the system can semantically match queries with relevant document chunks. Poor retrieval due to weak embeddings will degrade the overall

TABLE III
SYNTHETIC QA DATASET GENERATION PERFORMANCE ACROSS DOMAINS

Domain	Metric	Human	Ours	RAGAS
COVID-QA	Faithfulness	0.894	0.989	0.917
	Answerability	0.620	0.994	0.418
	Answer Relevance	0.399	0.947	0.748
RepLIQA	Faithfulness	0.774	0.994	0.957
	Answerability	0.740	0.998	0.916
	Answer Relevance	0.475	0.962	0.770
Army	Faithfulness	—	0.977	0.868
	Answerability	—	0.994	0.618
	Answer Relevance	—	0.957	0.830
Cybersecurity	Faithfulness	—	0.991	0.797
	Answerability	—	0.974	0.306
	Answer Relevance	—	0.961	0.783
Engineering	Faithfulness	—	0.988	0.839
	Answerability	—	0.990	0.473
	Answer Relevance	—	0.958	0.797

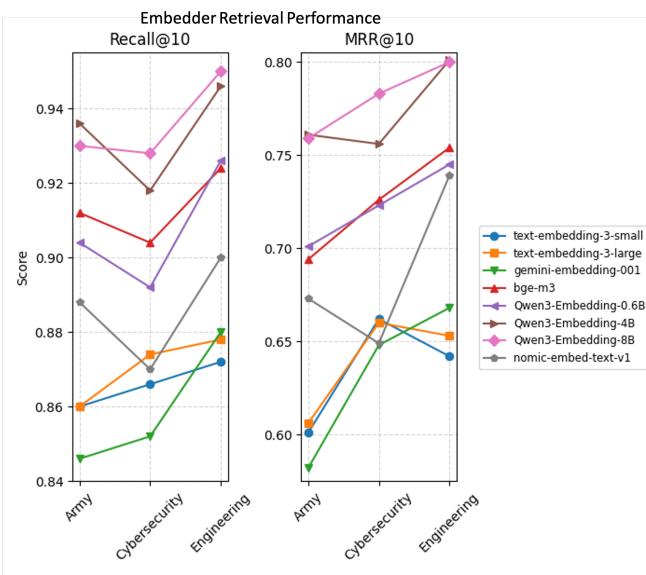


Fig. 3. We evaluate retrieval with Recall@10 and MRR@10 metrics on a variety of embedding models on three different domains.

output. Therefore, selecting the right model is a critical design decision. In this hypothetical scenario, we will evaluate the retrieval performance of a diverse selection of embedding models. Our analysis will include models that rank highly on the MTEB leaderboard [45], as well as a variety of open-source and closed-source models. We will also consider models of different sizes to assess the impact of model scale on performance.

To evaluate retrieval performance, we employ Recall@K and MRR@K on the datasets generated by our pipeline introduced in Section IV-B, using $k = 10$. Each document is chunked with a maximum size of 800 tokens and an overlap of 400 tokens between consecutive chunks. As shown in Figure 3, the results reveal the variation in embedding

model performance across different domains. The Qwen3 family of models [46] consistently performs well in retrieval tasks across various domains, reaffirming their high placement on the MTEB leaderboard. Despite its third-place ranking on MTEB, the gemini-embedding-001 [47] model underperforms compared to other embedding models on these specific domains. Meanwhile, the BGE-M3 [48] embedding model shows performance comparable to the Qwen3-Embedding-0.6B model, even though its MTEB ranking is significantly lower. Notably, open-source models generally outperform their closed-source counterparts. An interesting observation within the Qwen3 family is that the smaller 4B model performs similarly to, and in some domains even surpasses, its larger 8B counterpart.

Performance across different domains is inconsistent for most embedding models. For example, recall scores for most models are weaker in the cybersecurity domain compared to other areas. However, the text-embedding-3 family of embedding models defies this trend, performing much better in cybersecurity relative to the other domains. This suggests they were likely trained on a larger volume of cybersecurity-specific data. A similar pattern is observed with mean reciprocal rank (MRR), underscoring that embedding model effectiveness is highly domain-dependent. This inconsistency reinforces the need to evaluate RAG systems with domain-specific benchmarks.

B. Domain Specific LLM Generation Evaluation

This experiment evaluates the domain specific RAG performance of a diverse set of LLMs on Answer Correctness, Faithfulness, and Answer Relevancy. To ensure consistency, each model receives the same set of 10 retrieved context chunks with a maximum chunks size of 800 tokens, and chunk overlap of 400 tokens. For retrieval, we use the top-performing model on the MTEB retrieval task, Qwen3-Embedding-8B.

As shown in Figure 4, Gemini-2.5-flash [49] shows the strongest overall performance in Answer Correctness, achieving the highest scores in cybersecurity and bridge engineering domains. Most models struggle with the cybersecurity domain, with several models (including GPT-4o-mini, GPT-4.1, and Qwen3) dropping significantly in this domain compared to army and bridge engineering domains. This may be due to the drop in recall for Qwen3-Embedding-8B embedding model, as seen in Figure 3. Additionally, the Google family of models (Gemini and Gemma3 [50]) perform better than other models.

Gemini-2.5-flash demonstrates the strongest overall Faithfulness, achieving top scores in the army and cybersecurity and remaining competitive in bridge engineering domains. Models generally perform better in cybersecurity on Faithfulness, suggesting that the LLMs may not know the answer to the question from parametric knowledge and are leaning more heavily on the retrieved context chunks.

GPT-4.1-nano performs the best in all domains in Answer Relevancy. In general, the GPT family of models score higher than the Gemini models in this metric. This could be due to the fact that GPT models are more wordy in their responses,

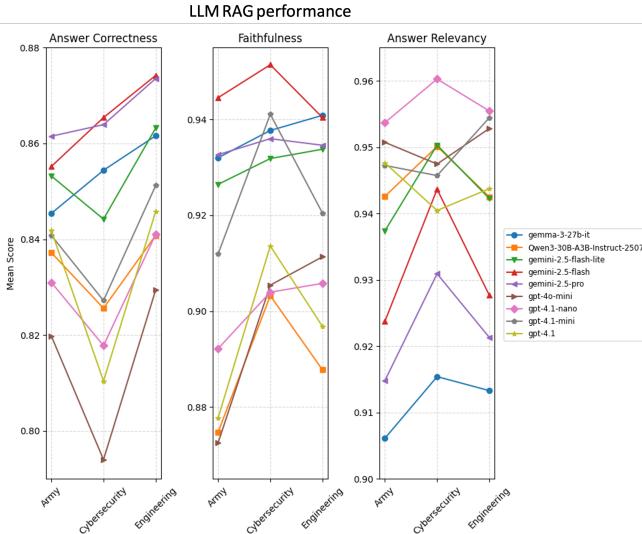


Fig. 4. We evaluate LLM generation with the Answer Correctness, Faithfulness, and Answer Relevancy metrics on three different domains.

making it more likely for the responses to be more relevant to the question.

No single model dominates all 3 metrics in all 3 domains. There doesn't seem to be a significant advantage of closed-sourced LLMs compared to open-sourced ones (Qwen3 and Gemma3) in these domains. Additionally, larger model sizes do not equate to better performance. This further highlighting the variability in performance across LLMs for different domains and the need for RAGalyst.

C. Domain Specific Retrieved Chunks Evaluation

To study the sensitivity of evaluation metrics to the number of retrieved chunks, we varied the number of chunks from 1 to 10 using Gemma3-4B. Across the three domains, the optimal number of chunks for maximizing Answer Correctness differed, suggesting that retrieval depth should be tuned to the domain. Moreover, all metrics exhibited substantial variation across domains, underscoring the importance of domain-specific RAG evaluation rather than a one-size-fits-all approach.

Despite these differences, some consistent trends emerge across metrics. Faithfulness tends to decline slightly as the number of chunks increases, likely because the LLM must contend with more irrelevant information. In contrast, Answer Relevancy generally improves with additional chunks, as the broader context encourages the model to at least address the question. Answer Correctness shows a peaked behavior, typically highest when 3–5 chunks are retrieved. With too few chunks, relevant information is often missed, reducing correctness. With too many, the relevant chunk risks being diluted by irrelevant ones, leading to distraction and lower performance.

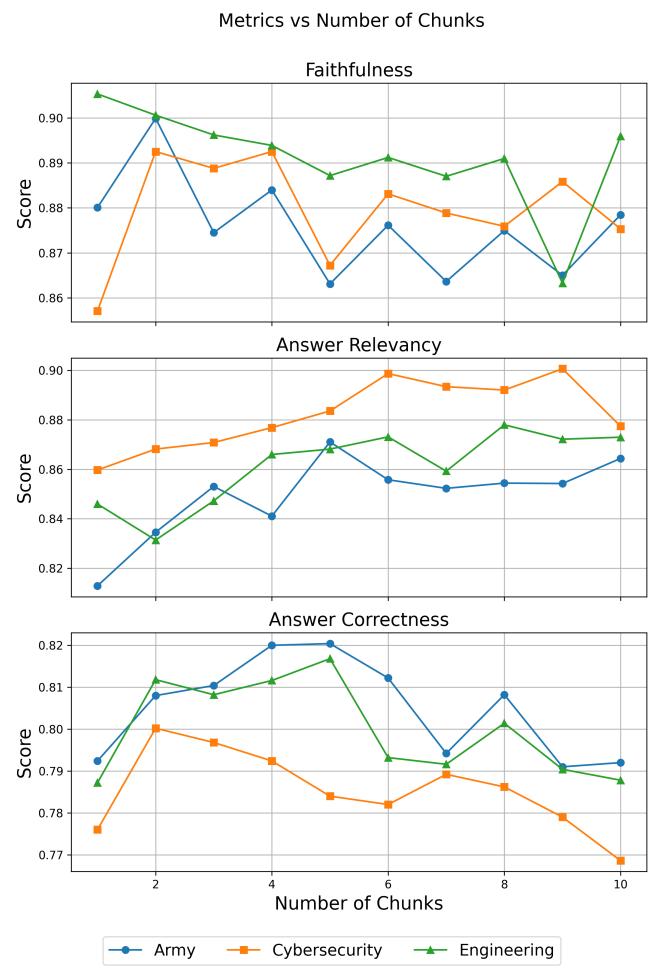


Fig. 5. We ablate the number of chunks retrieved with Gemma3-4B to assess the effect on LLM generation performance on Answer Correctness, Faithfulness and Answer Relevancy. This figure shows the each metric responds differently to the number of chunks retrieved, and that the ideal number of chunks retrieved to maximize Answer Correctness will vary.

D. Low Answer Correctness Analysis

Despite the promising performance of RAG, Answer Correctness remains imperfect, as shown with the non-perfect Answer Correctness in Section V-B. To diagnose the underlying causes, we employ GPT-5 as an LLM-as-a-Judge and analyze failures using a combined taxonomy derived from Barnett et al. [51] for RAG systems and Huang et al. [52] for LLMs. We exclude Missing Content from the RAG taxonomy because our agentic dataset generation pipeline ensures that all QAs are grounded in the source documents. Likewise, we remove Incomplete since the pipeline does not generate multi-part questions. Additionally, we refine Incorrect Specificity into two subcategories—Over Specificity and Under Specificity—to better capture granularity-related errors. Each QA may exhibit multiple failure types.

To assess GPT-5's reliability as an LLM-as-a-Judge, we manually validate its evaluations on ten QAs from each of the three domain-specific datasets. Across all thirty samples,

TABLE IV
LOW CORRECTNESS REASONS COUNTS AS A PERCENTAGE OF TOTAL QAS ACROSS THREE DOMAINS (MULTIPLE REASONS PER QA POSSIBLE)

All QAs				
Failure Types	Army	Cyber.	Engine.	Total
Number of QAs	500	500	500	1500
RAG Failures				
Over-Specificity	58.4%	81.0%	74.6%	71.3%
Not Extracted	11.0%	13.0%	15.0%	13.0%
Under-Specificity	9.8%	10.2%	10.4%	10.13%
Missed Top Ranked	5.0%	6.6%	5.4%	5.7%
Wrong Format	2.2%	0.6%	1.8%	1.5%
LLM Failures				
Context Inconsistency	6.6%	9.6%	8.8%	8.3%
Factual Fabrication	2.6%	7.4%	8.4%	6.1%
Factual Contradiction	2.6%	5.6%	3.6%	3.9%
Logical Inconsistency	0.8%	1.0%	0.4%	0.7%
Instruction Inconsistency	0.4%	0.0%	0.0%	0.1%
No Failures	31.1%	10.6%	14.4%	18.7%

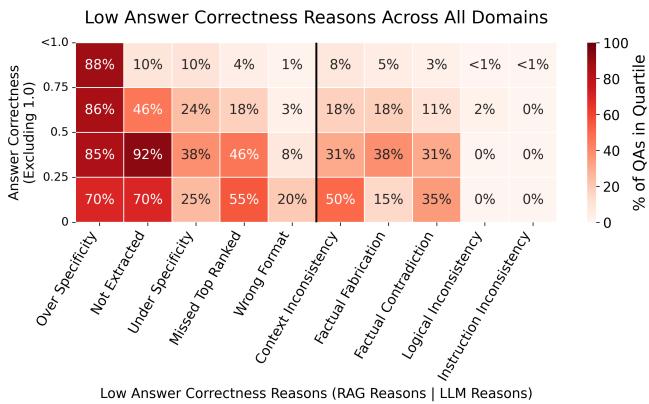


Fig. 6. Using GPT-5, we analyze the underlying reasons why Answer Correctness is low by quartile excluding 1.0. We show the number of low Answer Correctness reason as a percentage of the total number of QAs in each quartile. Note that percentages do not sum to 100% since each QA may exhibit multiple failure reasons. Reasons are grouped by taxonomy with RAG reasons on the left, and LLM reasons on the right.

GPT-5 consistently identifies all correct underlying reasons for low Answer Correctness.

As summarized in Table IV, Over Specificity accounts for the majority of failures. This pattern arises because QA generation relies on a single text chunk as context, leading GPT-4o-mini to produce concise answers constrained by limited information. In contrast, during RAG QA, GPT-4o-mini has access to a larger range retrieved contexts and tends to generate more verbose answers that incorporate all available information.

For answers scoring above 0.75 in Answer Correctness, as shown in 6, Over Specificity accounts for the vast majority of correctness issues. As Answer Correctness decline, QAs increasingly exhibit multiple contributing factors that lower their overall correctness. More importantly, all other failure reasons still exist, highlighting the weakness in RAG systems.

TABLE V
LLM PREFERENCE FOR SELF-GENERATED DATASETS

Evaluated Model	Metric	Dataset Origin	Score
GPT-4o-mini	Answer Correctness	GPT-4o-mini	0.820
		Gemini-2.5-flash	0.852
		Qwen3-30B...	0.827
	Faithfulness	GPT-4o-mini	0.873
		Gemini-2.5-flash	0.859
		Qwen3-30B...	0.858
Gemini-2.5-flash	Answer Relevancy	GPT-4o-mini	0.951
		Gemini-2.5-flash	0.930
		Qwen3-30B...	0.960
	Answer Correctness	GPT-4o-mini	0.852
		Gemini-2.5-flash	0.874
		Qwen3-30B...	0.897
Qwen3-30B...	Faithfulness	GPT-4o-mini	0.859
		Gemini-2.5-flash	0.942
		Qwen3-30B...	0.926
	Answer Relevancy	GPT-4o-mini	0.930
		Gemini-2.5-flash	0.921
		Qwen3-30B...	0.909
Qwen3-30B...	Answer Correctness	GPT-4o-mini	0.876
		Gemini-2.5-flash	0.859
		Qwen3-30B...	0.888
	Faithfulness	GPT-4o-mini	0.888
		Gemini-2.5-flash	0.879
		Qwen3-30B...	0.877
Qwen3-30B...	Answer Relevancy	GPT-4o-mini	0.929
		Gemini-2.5-flash	0.940
		Qwen3-30B...	0.932

E. Bias from Dataset Generation LLM

Recent studies suggest that LLMs often perform better on text they themselves generate [53]. To investigate whether this phenomenon introduces bias in RAG performance from our dataset generation pipeline, we construct datasets of 500 questions each using GPT-4o-mini, Gemini-2.5-flash, and Qwen3-30B-A3B-Instruct-2507 within the Army domain. Each dataset is then evaluated with all three models. Retrieval is performed with Qwen3-Embedding-8B, using a chunk size of 800 tokens, 400 token overlap, and top 10 chunk retrieval.

As shown in Table 4, we observe minimal evidence of bias from dataset origin. The only instance where Answer Correctness is highest on its own dataset occurs with Qwen3-30B-A3B-Instruct-2507. For Faithfulness, only Gemini-2.5-flash achieves its best score on its own dataset. For Answer Relevancy, none of the models perform best on their respective dataset origins.

VI. CONCLUSION

We present RAGalyst, an automated human-aligned agentic framework for domain-specific RAG evaluation. We achieve strong human alignment through prompt optimization for Answer Correctness and Answerability. Leveraging these metrics, our framework generates high-quality synthetic QA datasets that outperform both handcrafted benchmarks and RAGAS

across all metrics, enabling reliable evaluation of retrieval and generation without human supervision.

Our experiments reveal that RAG performance is highly configuration-dependent, with no universally optimal setup. Embedding model performance varies substantially across domains, often contradicting MTEB rankings. LLM generation performance differs across model families, with closed-source models showing no consistent advantage over open-source alternatives and larger models not consistently outperforming smaller ones. Our low Answer Correctness analysis reveals that RAG systems exhibit imperfect performance, with Incorrect Specificity, Incomplete Extraction, and Context Inconsistency emerging as the three most pressing failure modes.

These findings demonstrate the necessity of a systematic evaluation framework like RAGalyst, which enables practitioners to uncover domain-specific trade-offs and make informed design choices for building reliable RAG systems.

ACKNOWLEDGMENTS

Research was sponsored by the DEVCOM Analysis Center and was accomplished under Contract Number W911QX-23-D-0009. The views and conclusions contained in this document are those of the authors and should not be interpreted as representing the official policies, either expressed or implied, of the DEVCOM Analysis Center or the U.S. Government. The U.S. Government is authorized to reproduce and distribute reprints for Government purposes notwithstanding any copyright notation herein.

REFERENCES

- [1] A. T. Kalai, O. Nachum, S. S. Vempala, and E. Zhang, “Why language models hallucinate,” *arXiv preprint arXiv:2509.04664*, 2025.
- [2] L. Huang, W. Yu, W. Ma, W. Zhong, Z. Feng, H. Wang, Q. Chen, W. Peng, X. Feng, B. Qin *et al.*, “A survey on hallucination in large language models: Principles, taxonomy, challenges, and open questions,” *ACM Transactions on Information Systems*, vol. 43, no. 2, pp. 1–55, 2025.
- [3] N. Varshney, W. Yao, H. Zhang, J. Chen, and D. Yu, “A stitch in time saves nine: Detecting and mitigating hallucinations of llms by validating low-confidence generation,” *arXiv preprint arXiv:2307.03987*, 2023.
- [4] J. Li, J. Chen, R. Ren, X. Cheng, W. X. Zhao, J.-Y. Nie, and J.-R. Wen, “The dawn after the dark: An empirical study on factuality hallucination in large language models,” *arXiv preprint arXiv:2401.03205*, 2024.
- [5] Y. Gao, Y. Xiong, X. Gao, K. Jia, J. Pan, Y. Bi, Y. Dai, J. Sun, M. Wang, and H. Wang, “Retrieval-augmented generation for large language models: A survey,” *arXiv preprint arXiv:2312.10997*, 2023.
- [6] J. Li, Y. Yuan, and Z. Zhang, “Enhancing llm factual accuracy with rag to counter hallucinations: A case study on domain-specific queries in private knowledge-bases,” *arXiv preprint arXiv:2403.10446*, 2024.
- [7] G. Izacard, P. Lewis, M. Lomeli, L. Hosseini, F. Petroni, T. Schick, J. Dwivedi-Yu, A. Joulin, S. Riedel, and E. Grave, “Atlas: Few-shot learning with retrieval augmented language models,” *arXiv preprint arXiv:2208.03299*, 2022.
- [8] S. Es, J. James, L. Espinosa-Anke, and S. Schockaert, “Ragas: Automated evaluation of retrieval augmented generation,” *arXiv preprint arXiv:2309.15217*, 2023.
- [9] K. Papineni, S. Roukos, T. Ward, and W.-J. Zhu, “Bleu: a method for automatic evaluation of machine translation,” in *Proceedings of the 40th annual meeting of the Association for Computational Linguistics*, 2002, pp. 311–318.
- [10] C.-Y. Lin, “Rouge: a package for automatic evaluation of summaries,” in *Workshop on Text Summarization Branches Out, Post-Conference Workshop of ACL 2004, Barcelona, Spain*, July 2004, pp. 74–81. [Online]. Available: <https://www.microsoft.com/en-us/research/publication/rouge-a-package-for-automatic-evaluation-of-summaries/>
- [11] K. Zhu, Y. Luo, D. Xu, Y. Yan, Z. Liu, S. Yu, R. Wang, S. Wang, Y. Li, N. Zhang, X. Han, Z. Liu, and M. Sun, “Rageval: Scenario specific rag evaluation dataset generation framework,” *arXiv preprint arXiv:2408.01262*, 2025.
- [12] J. Saad-Falcon, O. Khattab, C. Potts, and M. Zaharia, “Ares: An automated evaluation framework for retrieval-augmented generation systems,” *arXiv preprint arXiv:2311.09476*, 2023.
- [13] Y. Liu, L. Huang, S. Li, S. Chen, H. Zhou, F. Meng, J. Zhou, and X. Sun, “Recall: A benchmark for llms robustness against external counterfactual knowledge,” *arXiv preprint arXiv:2311.08147*, 2023.
- [14] J. Chen, H. Lin, X. Han, and L. Sun, “Benchmarking large language models in retrieval-augmented generation,” in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 38, no. 16, 2024, pp. 17754–17762.
- [15] X. Yu, H. Cheng, X. Liu, D. Roth, and J. Gao, “Reeval: Automatic hallucination evaluation for retrieval-augmented large language models via transferable adversarial attacks,” *arXiv preprint arXiv:2310.12516*, 2023.
- [16] P. Lewis, E. Perez, A. Piktus, F. Petroni, V. Karpukhin, N. Goyal, H. Küttler, M. Lewis, W. tau Yih, T. Rocktäschel, S. Riedel, and D. Kiela, “Retrieval-augmented generation for knowledge-intensive nlp tasks,” 2021. [Online]. Available: <https://arxiv.org/abs/2005.11401>
- [17] Y. Gao, Y. Xiong, M. Wang, and H. Wang, “Modular rag: Transforming rag systems into lego-like reconfigurable frameworks,” *arXiv preprint arXiv:2407.21059*, 2024.
- [18] S. Jeong, J. Baek, S. Cho, S. J. Hwang, and J. C. Park, “Adaptive-rag: Learning to adapt retrieval-augmented large language models through question complexity,” *arXiv preprint arXiv:2403.14403*, 2024.
- [19] D. Edge, H. Trinh, N. Cheng, J. Bradley, A. Chao, A. Mody, S. Truitt, D. Metropolitansky, R. O. Ness, and J. Larson, “From local to global: A graphrag approach to query-focused summarization,” *arXiv preprint arXiv:2404.16130*, 2024.
- [20] Z. Jiang, X. Ma, and W. Chen, “Longrag: Enhancing retrieval-augmented generation with long-context llms,” *arXiv preprint arXiv:2406.15319*, 2024.
- [21] T. Yu, A. Xu, and R. Akkiraju, “In defense of rag in the era of long-context language models,” *arXiv preprint arXiv:2409.01666*, 2024.
- [22] T. Zhang, S. G. Patil, N. Jain, S. Shen, M. Zaharia, I. Stoica, and J. E. Gonzalez, “Raft: Adapting language model to domain specific rag,” *arXiv preprint arXiv:2403.10131*, 2024.
- [23] Z. Nguyen, A. Annunziata, V. Luong, S. Dinh, Q. Le, A. H. Ha, C. Le, H. A. Phan, S. Raghavan, and C. Nguyen, “Enhancing q&a with domain-specific fine-tuning and iterative reasoning: A comparative study,” *arXiv preprint arXiv:2404.11792*, 2024.
- [24] S. Wang, J. Liu, S. Song, J. Cheng, Y. Fu, P. Guo, K. Fang, Y. Zhu, and Z. Dou, “Domainrag: A chinese benchmark for evaluating domain-specific retrieval-augmented generation,” *arXiv preprint arXiv:2406.05654*, 2024.
- [25] H. Yu, A. Gan, K. Zhang, S. Tong, Q. Liu, and Z. Liu, “Evaluation of retrieval-augmented generation: A survey,” *arXiv preprint arXiv:2405.07437*, 2024.
- [26] A. Salemi and H. Zamani, “Evaluating retrieval quality in retrieval-augmented generation,” *arXiv preprint arXiv:2404.13781*, 2024.
- [27] J. Liu, R. Ding, L. Zhang, P. Xie, and F. Huang, “Cofe-rag: A comprehensive full-chain evaluation framework for retrieval-augmented generation with enhanced data diversity,” *arXiv preprint arXiv:2410.12248*, 2024.
- [28] Y. Liu, T. Yang, S. Huang, Z. Zhang, H. Huang, F. Wei, W. Deng, F. Sun, and Q. Zhang, “Calibrating llm-based evaluator,” *arXiv preprint arXiv:2309.13308*, 2023.
- [29] O. Khattab, A. Singhvi, P. Maheshwari, Z. Zhang, K. Santhanam, S. Vardhaman, S. Haq, A. Sharma, T. T. Joshi, H. Moazam, H. Miller, M. Zaharia, and C. Potts, “Dspy: Compiling declarative language model calls into self-improving pipelines,” 2024.
- [30] O. Khattab, K. Santhanam, X. L. Li, D. Hall, P. Liang, C. Potts, and M. Zaharia, “Demonstrate-search-predict: Composing retrieval and language models for knowledge-intensive NLP,” *arXiv preprint arXiv:2212.14024*, 2022.
- [31] P. Verga, S. Hofstätter, S. Althammer, Y. Su, A. Piktus, A. Arkhangorodsky, M. Xu, N. White, and P. Lewis, “Replacing judges with juries: Evaluating llm generations with a panel of diverse models,” *arXiv preprint arXiv:2404.18796*, 2024.

- [32] F. Bai, K. Harrigan, J. Stremmel, H. Hassanzadeh, A. Saeedi, and M. Dredze, "Give me some hard questions: Synthetic data generation for clinical qa," *arXiv preprint arXiv:2412.04573*, 2024.
- [33] C. Alberti, D. Andor, E. Pitler, J. Devlin, and M. Collins, "Synthetic qa corpora generation with roundtrip consistency," *arXiv preprint arXiv:1906.05416*, 2019.
- [34] S. Shakeri, C. N. dos Santos, H. Zhu, P. Ng, F. Nan, Z. Wang, R. Nallapati, and B. Xiang, "End-to-end synthetic data generation for domain adaptation of question answering systems," *arXiv preprint arXiv:2010.06028*, 2020.
- [35] B. Bohnet, K. Swersky, R. Liu, P. Awasthi, A. Nova, J. Snaider, H. Sedghi, A. Parisi, M. Collins, A. Lazaridou, O. Firat, and N. Fiedel, "Long-span question-answering: Automatic question generation and qa-system ranking via side-by-side evaluation," *arXiv preprint arXiv:2406.00179*, 2024.
- [36] A. Jadon, A. Patil, and S. Kumar, "Enhancing domain-specific retrieval-augmented generation: Synthetic data generation and evaluation using reasoning models," *arXiv preprint arXiv:2502.15854*, 2025.
- [37] J. Gu, Z. S. Xuhui Jiang, X. Z. Hexiang Tan, W. L. Chengjin Xu, S. M. Yinghan Shen, S. W. Honghao Liu, Y. W. Kun Zhang, L. N. Wen Gao, and J. Guo, "A survey on llm-as-a-judge," *arXiv preprint arXiv:2411.15594*, 2025.
- [38] Y. Liu, Y. X. Dan Iter, R. X. Shuohang Wang, and C. Zhu, "G-eval: NLG evaluation without human-label bias," *arXiv preprint arXiv:2305.14228*, 2023.
- [39] L. Zheng, Y. S. Wei-Lin Chiang, Z. W. Siyuan Zhuang, Z. L. Yong-hao Zhuang, D. L. Zhuohan Li, H. Z. E. Xing, J. E. Gonzalez, and I. Stoica, "Judging llm-as-a-judge with mt-bench and chatbot arena," in *Proceedings of NeurIPS*, 2023.
- [40] T. Kocmi and C. Federmann, "Large language models are state-of-the-art evaluators of translation quality," *arXiv preprint arXiv:2304.13756*, 2023.
- [41] D. G. Bonett and T. A. Wright, "Sample size requirements for estimating pearson, kendall and spearman correlations," *Psychometrika*, vol. 65, no. 1, pp. 23–28, 2000.
- [42] T. Möller, A. Reina, R. Jayakumar, and M. Pietsch, "COVID-QA: A question answering dataset for COVID-19," in *Proceedings of the NLP COVID-19 Workshop at ACL 2020*, 2020. [Online]. Available: <https://aclanthology.org/2020.nlpcovid19-acl.18>
- [43] C. Pal, Z. Jin, T. Wu, S. Abnar, E. Voorhees, and S. Bengio, "Replqa: A benchmark for retrieving and reading unseen fictional documents," in *Proceedings of the TMLR Conference*, 2024, accessed: 2025-06-25. [Online]. Available: <https://openreview.net/forum?id=4diKTLmg2y>
- [44] M. Gozzi and F. D. Maio, "Comparative analysis of prompt strategies for large language models: Single-task vs. multitask prompts," *Electronics*, vol. 13, no. 23, p. 4712, 2024. [Online]. Available: <https://doi.org/10.3390/electronics13234712>
- [45] N. Muennighoff, N. Tazi, L. Magne, and N. Reimers, "Mteb: Massive text embedding benchmark," *arXiv preprint arXiv:2210.07316*, 2022.
- [46] Y. Zhang, M. Li, D. Long, X. Zhang, H. Lin, B. Yang, P. Xie, A. Yang, D. Liu, J. Lin *et al.*, "Qwen3 embedding: Advancing text embedding and reranking through foundation models," *arXiv preprint arXiv:2506.05176*, 2025.
- [47] J. Lee, F. Chen, S. Dua, D. Cer, M. Shanbhogue, I. Naim, G. H. Ábrego, Z. Li, K. Chen, H. S. Vera *et al.*, "Gemini embedding: Generalizable embeddings from gemini," *arXiv preprint arXiv:2503.07891*, 2025.
- [48] J. Chen, S. Xiao, P. Zhang, K. Luo, D. Lian, and Z. Liu, "M3-embedding: Multi-linguality, multi-functionality, multi-granularity text embeddings through self-knowledge distillation," in *Findings of the Association for Computational Linguistics ACL 2024*, 2024, pp. 2318–2335.
- [49] G. Comanici, E. Bieber, M. Schaeckermann, I. Pasupat, N. Sachdeva, I. Dhillon, M. Blstein, O. Ram, D. Zhang, E. Rosen *et al.*, "Gemini 2.5: Pushing the frontier with advanced reasoning, multimodality, long context, and next generation agentic capabilities," *arXiv preprint arXiv:2507.06261*, 2025.
- [50] G. Team, A. Kamath, J. Ferret, S. Pathak, N. Vieillard, R. Merhej, S. Perrin, T. Matejovicova, A. Ramé, M. Rivière *et al.*, "Gemma 3 technical report," *arXiv preprint arXiv:2503.19786*, 2025.
- [51] S. Barnett, S. Kurniawan, S. Thudumu, Z. Brannelly, and M. Abdelrazeq, "Seven failure points when engineering a retrieval augmented generation system," 2024. [Online]. Available: <https://arxiv.org/abs/2401.05856>
- [52] L. Huang, W. Yu, W. Ma, W. Zhong, Z. Feng, H. Wang, Q. Chen, W. Peng, X. Feng, B. Qin, and T. Liu, "A survey on hallucination in large language models: Principles, taxonomy, challenges, and open questions," *ACM Transactions on Information Systems*, vol. 43, no. 2, p. 1–55, Jan. 2025. [Online]. Available: <http://dx.doi.org/10.1145/3703155>
- [53] W. Xu, G. Zhu, X. Zhao, L. Pan, L. Li, and W. Y. Wang, "Pride and prejudice: Llm amplifies self-bias in self-refinement," *arXiv preprint arXiv:2402.11436*, 2024.

VII. SUPPLEMENTARY MATERIAL

A. Experimental Figures

We provide the exact figures from our experimental section V. We abbreviate the Military Operations domain to Army, Cybersecurity domain to Cyber, and the the Bridge Engineering domain to Eng.

1) *Embedding Retrieval Evaluation Across Domains*: The performance results for each embedding model and domain from Figure 3 are reported in Table VI.

TABLE VI
PERFORMANCE COMPARISON OF TEXT EMBEDDING MODELS ACROSS DOMAINS.

Model	Recall@10			MRR@10		
	Army	Cyber	Eng	Army	Cyber	Eng
text-embedding-3-small	0.864	0.866	0.872	0.601	0.662	0.641
text-embedding-3-large	0.860	0.874	0.878	0.606	0.660	0.653
gemini-embedding-001	0.846	0.852	0.880	0.582	0.648	0.668
bge-m3	0.912	0.904	0.924	0.694	0.726	0.754
Qwen3-Embedding-0.6B	0.904	0.892	0.926	0.701	0.723	0.745
Qwen3-Embedding-4B	0.936	0.918	0.946	0.761	0.756	0.801
Qwen3-Embedding-8B	0.930	0.928	0.950	0.759	0.783	0.800
nomic-embed-text-v1	0.888	0.870	0.900	0.673	0.649	0.739

2) *Domain Specific LLM Generation Evaluation*: The performance results for each LLM and domain from Figure 4 are reported in Table VII.

TABLE VII
PERFORMANCE OF LLMs ACROSS DOMAINS FOR ANSWER CORRECTNESS, FAITHFULNESS, AND ANSWER RELEVANCY.

Model	Answer Correctness			Faithfulness			Answer Relevancy		
	Army	Cyber	Eng	Army	Cyber	Eng	Army	Cyber	Eng
gemma-3-27b-it	0.86	0.84	0.86	0.93	0.94	0.93	0.91	0.91	0.92
Qwen3-30B-A3B-Instruct-2507	0.84	0.82	0.84	0.88	0.87	0.90	0.93	0.94	0.93
gemini-2.5-flash-lite	0.85	0.84	0.86	0.93	0.92	0.94	0.94	0.96	0.94
gemini-2.5-flash	0.86	0.85	0.88	0.95	0.94	0.95	0.95	0.96	0.95
gemini-2.5-pro	0.87	0.84	0.88	0.94	0.93	0.95	0.95	0.95	0.96
gpt-4.0-mini	0.80	0.79	0.82	0.88	0.90	0.91	0.94	0.95	0.94
gpt-4.1-nano	0.83	0.81	0.85	0.89	0.90	0.92	0.93	0.93	0.95
gpt-4.1-mini	0.85	0.84	0.86	0.91	0.94	0.93	0.92	0.94	0.94
gpt-4.1	0.84	0.82	0.86	0.92	0.91	0.93	0.94	0.95	0.95

3) *Domain Specific Retrieved Chunks Evaluation*: The performance results of Gemma3-4B on Answer Correctness, Faithfulness, and Answer Relevancy on all domains from Figure 5 are reported in Table VIII

TABLE VIII
PERFORMANCE OF DOMAINS ACROSS VARYING NUMBERS OF CHUNKS.

Metric	Domain	1	2	3	4	5	6	7	8	9	10
Faithfulness	Army	0.880	0.900	0.874	0.884	0.863	0.876	0.864	0.875	0.865	0.878
	Cyber	0.857	0.893	0.889	0.893	0.867	0.883	0.879	0.876	0.886	0.875
	Eng	0.905	0.901	0.896	0.894	0.887	0.891	0.887	0.891	0.863	0.896
Answer Relevancy	Army	0.813	0.835	0.853	0.841	0.871	0.856	0.852	0.854	0.854	0.864
	Cyber	0.860	0.868	0.871	0.877	0.884	0.899	0.893	0.892	0.901	0.877
	Eng	0.846	0.831	0.847	0.866	0.868	0.873	0.859	0.878	0.872	0.873
Answer Correctness	Army	0.792	0.808	0.810	0.820	0.820	0.812	0.794	0.808	0.791	0.792
	Cyber	0.776	0.800	0.797	0.792	0.784	0.782	0.789	0.786	0.779	0.769
	Eng	0.787	0.812	0.808	0.812	0.817	0.793	0.792	0.801	0.790	0.788

B. LLM Prompts

1) *Hand-Crafted Answer Correctness Prompt:* The following is our hand-crafted prompt for Answer Correctness.

Evaluation Prompt

You will be given a student answer and a ground truth.

Your task is to evaluate the student answer by comparing it with the ground truth. Give your evaluation on a scale of 0.0 to 1.0, where 0.0 means that the answer is completely unrelated to the ground truth, and 1.0 means that the answer is completely accurate and aligns perfectly with the ground truth.

For instance,

correctness_score: 0.0 – The answer is completely unrelated to the ground truth.

correctness_score: 0.3 – The answer has minor relevance but does not align with the ground truth.

correctness_score: 0.5 – The answer has moderate relevance but contains inaccuracies.

correctness_score: 0.7 – The answer aligns with the reference but has minor errors or omissions.

correctness_score: 1.0 – The answer is completely accurate and aligns perfectly with the ground truth.

You must provide values for correctness_score: in your answer.

Now here is the student answer and the ground truth.

2) *Prompt-Optimized Answer Correctness Prompt:* The following is the MIPROv2 and LabeledFewShot optimized Answer Correctness prompt.

Evaluation Prompt

```
{  
  "response": "8 rockets fired from Gaza into southern Israel; none hurt",  
  "reference": "Ten rockets from Gaza land in southern Israel; none hurt",  
  "correctness_score": 0.7  
},  
{  
  "response": "A person plays a keyboard.",  
  "reference": "Someone is playing a keyboard.",  
  "correctness_score": 1.0  
},  
{  
  "response": "What isn't how what was sold?",  
  "reference": "It's not how it was sold, gb.",  
  "correctness_score": 0.3  
},  
{  
  "response": "Jaya Prada all set to join BJP",  
  "reference": "Jaya Prada likely to join BJP, Amar Singh to decide for her",  
  "correctness_score": 0.8  
},  
{  
  "response": "Israel strikes Syria as tensions rise on weapons",  
  "reference": "Air strikes wound civilians in Syria's Deraa",  
  "correctness_score": 0.4  
},  
{  
  "response": "The issue has been resolved, Marlins President David Samson said through a club spokesman.",  
  "reference": "The Marlins only said: The issue has been resolved.",  
  "correctness_score": 0.6  
},  
{  
  "response": "Typhoon survivors raid Philippine stores",  
  "reference": "Typhoon Bopha kills 15 in S. Philippines",  
  "correctness_score": 0.2  
},  
{  
  "response": "three little boys cover themselves with bubbles.",  
  "reference": "Three children standing by a pool are covered in foam bubbles.",  
  "correctness_score": 0.8  
}
```

You are a language assessment evaluator. You will be given a student answer and a ground truth response. Your task is to evaluate the student answer by comparing it with the ground truth and provide a similarity score on a scale of 0.0 to 1.0. A score of 0.0 indicates that the answer is completely unrelated to the ground truth, while a score of 1.0 indicates that the answer is completely accurate and aligns perfectly with the ground truth. Please include the evaluation in the format: correctness_score: [score].

Now here is the student answer and the ground truth.

3) *Answerability*: The following is our hand-crafted prompt for Answerability.

Evaluation Prompt

You will be given a context and a question.

Your task is to determine if the question is clearly and unambiguously answerable using only the given context. - If the context contains **all** the necessary information to answer the question **without making assumptions** or using **any external knowledge**, then the groundedness is 1. - Otherwise, if any key information is **missing**, ambiguous, or requires inference beyond what is stated, then the groundedness is 0.

You **MUST** provide values for 'answerability_flag' in your answer.

Use only the provided context. Do not use prior knowledge, common sense, or information not explicitly contained in the context.