MEMORY-FREE CONTINUAL LEARNING WITH NULL SPACE ADAPTATION FOR ZERO-SHOT VISION-LANGUAGE MODELS

Yujin Jo and Taesup KimGraduate School of Data Science
Seoul National University

ABSTRACT

Pre-trained vision-language models (VLMs), such as CLIP, have demonstrated remarkable zero-shot generalization, enabling deployment in a wide range of realworld tasks without additional task-specific training. However, in real deployment scenarios with evolving environments or emerging classes, these models inevitably face distributional shifts and novel tasks. In such contexts, static zero-shot capabilities are insufficient, and there is a growing need for continual learning methods that allow models to adapt over time while avoiding catastrophic forgetting. We introduce NuSA-CL (Null Space Adaptation for Continual Learning), a lightweight memory-free continual learning framework designed to address this challenge. NuSA-CL employs low-rank adaptation and constrains task-specific weight updates to lie within an approximate null space of the model's current parameters. This strategy minimizes interference with previously acquired knowledge, effectively preserving the zero-shot capabilities of the original model. Unlike methods relying on replay buffers or costly distillation, NuSA-CL imposes minimal computational and memory overhead, making it practical for deployment in resource-constrained, real-world continual learning environments. Experiments show that our framework not only effectively preserves zero-shot transfer capabilities but also achieves highly competitive performance on continual learning benchmarks. These results position NuSA-CL as a practical and scalable solution for continually evolving zero-shot VLMs in real-world applications.

1 Introduction

Vision-language foundation models such as CLIP(Radford et al., 2021) have brought about a major shift in artificial intelligence by enabling zero-shot generalization. Their powerful text-image aligned representations now serve as the perceptual core for a new generation of systems, including Multimodal Large Language Models (MLLMs) like LLaVA(Liu et al., 2023; 2024) and Vision-Language Action (VLA) models for robotics(Kim et al., 2024; Shukor et al., 2025). However, these advanced systems inherit a critical limitation from their static backbones: In settings where data distributions and user requirements are constantly evolving, their knowledge is frozen.

To bridge the gap between static foundation models and the demands of real-world deployment without resorting to massive retraining, Continual Learning (CL) has emerged as a promising solution. CL allows models to incrementally acquire new knowledge while preventing catastrophic forgetting of both pre-trained and previously learned tasks. Existing CL paradigms, however, face a fundamental scalability wall. Storage-based methods, which rely on experience replay or reference data(Saha et al., 2021; Wang et al., 2021), are inherently constrained by storage costs that grow linearly with the number of tasks. On the other hand, expansion-based methods introduce new modules for each task, such as adapters or prompts(Yu et al., 2024; Tang et al., 2024), forcing a model's parameters and architectural complexity to grow unbounded over time. While effective on short-term benchmarks, these dominant approaches are ill-suited for true lifelong learning. Even many parameter-efficient fine-tuning (PEFT) techniques still depend on explicitly storing prior task information to mitigate interference(Liang & Li, 2024; Lu et al., 2024).

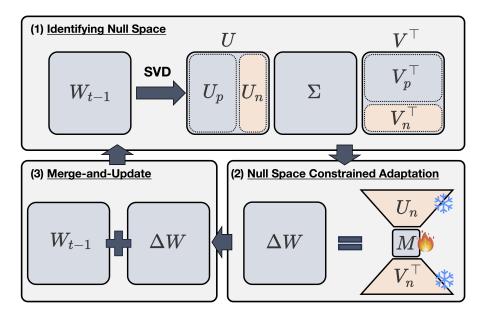


Figure 1: The NuSA-CL framework. Starting with the weights from the previous task W_{t-1} , we first perform SVD to identify the intrinsic null space. A new low-rank update ΔW_t is then learned under a persistent constraint that confines it to this space. Finally, the update is merged to produce the new weights $W_t \leftarrow W_{t-1} + \Delta W$, completing the cycle.

We argue that overcoming this scalability wall requires a paradigm shift from relying on external resources to enabling a model to adapt using only its intrinsic structure. We propose NuSA-CL (Null Space Adaptation for Continual Learning), a continual learning framework that enables a model with a fixed capacity to efficiently reorganize its own knowledge to accommodate new information. NuSA-CL dynamically identifies an underutilized null space in the model's current weights via SVD before each new task and strictly confines all weight updates within these interference-free dimensions throughout training.

This data-agnostic process concludes by merging the update into the backbone, maintaining a fixed parameter budget. By preserving the model's core knowledge, NuSA-CL enables stable continual adaptation, offering the ultimate form of scalability with zero storage overhead, zero auxiliary model load, and zero parameter growth which is a crucial set of properties for resource-constrained environments such as autonomous agents or on-device AI.

Our contributions are summarized as follows:

- We propose NuSA-CL, a novel memory-free and resource-efficient continual learning method for vision-language foundation models, designed to operate effectively in resource-constrained environments without relying on memory buffers or knowledge distillation.
- NuSA-CL introduces a null space-constrained low-rank update strategy that integrates new
 knowledge into an approximate null space of the pre-trained parameters, thereby preserving
 zero-shot generalization while enabling stable and incremental learning.
- Our method demonstrates significant computational and memory efficiency, making it well-suited for real-world applications that demand lifelong adaptability under limited resources.

2 RELATED WORK

2.1 CONTINUAL LEARNING WITH PARAMETER-EFFICIENT FINE-TUNING

Continual learning (CL) aims to adapt models to a sequence of tasks without the catastrophic forgetting of previously acquired knowledge (Li & Hoiem, 2017; Rebuffi et al., 2017; Chaudhry

et al., 2019; Ding et al., 2022). A central challenge is the *stability-plasticity trade-off*, which is amplified in foundation models like CLIP, where forgetting undermines not only past tasks but also the general-purpose zero-shot capabilities acquired during pre-training (Wortsman et al., 2022; Zhang et al., 2023; Tan et al., 2024). While full fine-tuning methods like ZSCL (Zheng et al., 2023) can be effective, they often require resource-intensive techniques, hindering their scalability.

Parameter-Efficient Fine-Tuning (PEFT) offers a lightweight alternative by restricting updates to a small subset of parameters. This family includes prompt-based methods that isolate task-specific knowledge (Wang et al., 2022b;a) and adapter-based methods that insert small, trainable modules for each task (Yu et al., 2024; Tang et al., 2024; Wu et al., 2025; Wei et al., 2024). However, many of these approaches still externalize new knowledge into task-specific modules, which contributes to parameter growth over long task sequences. In contrast, our work focuses on adapting the model's core weights within a fixed parameter budget.

2.2 ORTHOGONAL PROJECTION AND NULL SPACE APPROACHES

To explicitly mitigate interference, orthogonal projection techniques constrain parameter updates to subspaces that are orthogonal to those encoding prior knowledge. Prior works typically project new updates away from subspaces identified using stored information, such as past data, features, or gradients (Wang et al., 2021; Saha et al., 2021; Zhao et al., 2023). For instance, InfLoRA (Liang & Li, 2024) adapts this concept to LoRA but still necessitates a memory bank of past gradients to enforce orthogonality. This reliance on external memory contrasts with our strictly memory-free approach. Our method is distinct in that it derives the approximate null space intrinsically from the model's current weight structure via SVD, requiring no access to or storage of past data, features, or gradients.

2.3 SVD-GUIDED ADAPTATION IN FOUNDATION MODELS

Several recent methods explore using the spectral properties of a model's weights to guide adaptation, primarily for single-task fine-tuning (Lingam et al., 2024; Yang et al., 2025; Tang et al., 2025). These approaches involve adapting either principal components (Meng et al., 2024) or, conversely, minor low-energy components to minimize interference with pre-trained knowledge (Wang et al., 2025). While insightful, these methods differ from our work in two fundamental ways. First, they are designed for single-task adaptation, not the long-term, sequential learning required in CL. Second, and most critically, they use the low-energy subspace only for initialization, allowing the weight updates to deviate from this subspace during training. Our method, in contrast, enforces a persistent constraint, ensuring updates are strictly confined to the dynamically identified null space. This enables stable, lifelong learning within a fixed-capacity model.

3 METHOD: NULL SPACE ADAPTATION FOR CONTINUAL LEARNING

The core of NuSA-CL is a cyclical, data-agnostic adaptation process that enables a model to learn from a new task while preserving previously acquired knowledge. For each task in a sequence, the process unfolds in three stages as illustrated in Figure 1: (1) Null Space Identification via SVD: We begin with the model's current weights, W_{t-1} , and perform Singular Value Decomposition (SVD) to identify a low-energy subspace, that is, the intrinsic null space where prior knowledge is minimally encoded. (2) Constrained Adaptation: We then train a task-specific, low-rank update, ΔW_t , for the current task. Crucially, this update is persistently constrained to lie strictly within the identified null space throughout training. (3) Weight Merging: After training, the learned update is merged directly into the backbone weights, producing the updated model $W_t \leftarrow W_{t-1} + \Delta W_t$. This evolved model, with its fixed parameter budget, then serves as the starting point for the next task, where the cycle repeats.

3.1 IDENTIFYING THE INTRINSIC NULL SPACE

Let $W \in \mathbb{R}^{m \times n}$ be a weight matrix from the model. We compute its SVD, $W = U\Sigma V^{\top}$, to analyze its spectral properties. We posit that the principal components, associated with high-energy singular values, encode the core knowledge of the model. To provide a principled basis for our approach, we

first verify the existence of a sufficiently large low-energy subspace. We identify the dimension k of the principal space by finding the smallest integer that captures at least a ρ fraction of the total spectral energy:

$$\sum_{i=1}^{k} \sigma_i^2 \ge \rho \cdot \|W\|_F^2 \tag{1}$$

The remaining d-k dimensions constitute the intrinsic null space, spanned by the basis vectors (U_n, V_n) . For practical stability and to maintain a consistent number of trainable parameters across all layers and tasks, we cap the dimension of our update by a hyperparameter r_{max} . The effective rank r of our update is thus defined as $r = \min(d - k, r_{\text{max}})$.

We can now express the decomposition as:

$$U = [U_p \ U_n], \quad V = [V_p \ V_n], \quad \Sigma = \begin{bmatrix} \Sigma_p & 0 \\ 0 & \Sigma_n \end{bmatrix},$$
 (2)

where (U_p, V_p) represent the top-k components of the principal subspace, and (U_n, V_n) span the approximate null space.

3.2 Constrained Adaptation within the Null Space

To prevent interference with existing knowledge during continual learning, we impose a *persistent* constraint that strictly confines all new parameter updates to the identified null space. We formulate the task-specific adaptation as a LoRA-like low-rank update $\Delta W \in \mathbb{R}^{m \times n}$, but with a critical modification. Instead of learning two projection matrices, we define the update as:

$$\Delta W = U_n M V_n^{\top} \tag{3}$$

Here, the basis matrices U_n and V_n are derived from the SVD of the frozen weight W and are themselves kept frozen during training for the current task. The intermediate matrix $M \in \mathbb{R}^{r \times r}$ is the only trainable component and is initialized as a zero matrix for each new task. This formulation ensures that the update ΔW is mathematically guaranteed to be orthogonal to the principal subspace of W, thereby minimizing interference. This persistent constraint is a key distinction from prior work (Wang et al., 2025) that uses such subspaces only for initialization, after which updates are free to deviate.

3.3 CONTINUAL ADAPTATION VIA UPDATE MERGING

A core component of NuSA-CL's scalability is its ability to operate within a fixed parameter budget. This is achieved by merging the learned low-rank update ΔW directly into the base weights after training on each task is complete. For a given task t, the new weight matrix W_t is computed as:

$$W_t \leftarrow W_{t-1} + \Delta W_t \tag{4}$$

This update-and-merge cycle allows the model to sequentially accumulate knowledge from new tasks without adding any new parameters or modules. The resulting model, with its updated weights W_t , then becomes the starting point for the next task, t+1. At the beginning of the new task, the entire process repeats: the now-updated weights W_t are decomposed via SVD to identify a new intrinsic null space, ensuring that the model is always adapting in directions that are least disruptive to its full, accumulated knowledge.

4 THEORETICAL MOTIVATION

In this section, we provide a theoretical motivation for our approach. We analyze the degree of interference in *parameter space* to demonstrate how our persistent constraint provides a principled mechanism for mitigating catastrophic forgetting. Our analysis shows that by freezing the null space basis vectors (U_n, V_n) and only learning the small intermediate matrix M defined in Eq. 3, the update direction is guaranteed to be nearly orthogonal to the dominant components of the existing model weights.

4.1 Interference Bound for a Single Update

We first present a lemma that characterizes the interaction between the existing weights and a single task-specific update.

Lemma 1 (Bounded Interference via Null Space Constraint). Let $W = U\Sigma V^{\top}$ be the SVD of a weight matrix, and let $\Delta W = U_n M V_n^{\top}$ be an update restricted to its intrinsic null space. The interference in parameter space, measured by the Frobenius inner product, is bounded by:

$$|\langle W, \Delta W \rangle_F| \le \sigma_{max}^{null} \cdot ||M||_F \tag{5}$$

where $\sigma_{max}^{null} := \sigma_{k+1}$ is the largest singular value within the null space.

Proof. Expanding the inner product:
$$\langle W, \Delta W \rangle_F = \operatorname{Tr}(W^\top \Delta W) = \operatorname{Tr}(V \Sigma U^\top U_n M V_n^\top)$$
. Since $U^\top U_n = [0; I_r]$ and $V^\top V_n = [0; I_r]$, the trace simplifies to: $\operatorname{Tr}(\Sigma_n M) \leq \|\Sigma_n\|_2 \cdot \|M\|_F = \sigma_{\max}^{null} \cdot \|M\|_F$.

4.2 FORGETTING CONTROL IN CONTINUAL LEARNING

The above lemma naturally generalizes to the continual learning setting, where multiple tasks are learned sequentially.

Theorem 2 (Cumulative Interference Bound). Let $W_t = W_{t-1} + \Delta W_t$, where $\Delta W_t = U_{t-1,n} M_t V_{t-1,n}^{\top}$ is the update for task t. The cumulative interference across T tasks is bounded by:

$$\sum_{t=1}^{T} |\langle W_{t-1}, \Delta W_t \rangle_F| \le \sum_{t=1}^{T} \sigma_{t, max}^{null} \cdot ||M_t||_F.$$
 (6)

This result demonstrates that by constraining updates to low-energy subspaces, NuSA-CL bounds the cumulative parameter-level interference across tasks. This provides a principled mechanism for mitigating catastrophic forgetting, as it minimizes disruptions to the dominant weight structures that encode prior knowledge.

5 EXPERIMENTS

5.1 EXPERIMENTAL SETUP

Benchmarks. Our primary evaluation is conducted on the *Multimodal Task Incremental Learning (MTIL) benchmark* (Zheng et al., 2023), a sequence of 11 diverse vision datasets designed to test a model's ability to learn new tasks while preserving its core zero-shot capabilities. To assess long-sequence scalability, we also evaluate on the standard *Class-Incremental CIFAR100* benchmark (Krizhevsky et al., 2009), splitting 100 classes into sequences of 10, 20, and 50 tasks. We report three key metrics: *Transfer*, the zero-shot accuracy on unseen tasks; *Avg.*, the average accuracy across all tasks during training; and *Last*, the final average accuracy, which measures forgetting.

Implementation and Baselines. All experiments use the CLIP ViT-B/16 backbone. Our method, NuSA-CL, identifies the null space for each task using a cumulative energy cutoff and caps the LoRA update rank at $r_{\rm max}=128$. We compare NuSA-CL against three categories of baselines. (1) Full Fine-Tuning Models (e.g., Continual-FT, ZSCL) update all 150M parameters and require significantly more computational resources (e.g., 4 GPUs in our experiments) compared to the single GPU usage of PEFT methods. (2) Storage-based PEFT Models require additional, often expanding, storage, including MoE-Adapters (Yu et al., 2024), DIKI (Tang et al., 2024), and InfLoRA (Liang & Li, 2024). (3) Storage-Free PEFT Models, the most practical and challenging setting, includes standard LoRA (Hu et al., 2021), MiLoRA (Wang et al., 2025), and our method. For methodological relevance, we re-implemented the most comparable LoRA-based methods within a unified framework, applying adapters to both vision and text encoders with a consistent rank and merging them after each task.

Table 1: Performance and Computational Efficiency Analysis on the MTIL Benchmark. Boldface indicates the top storage-free performer. † indicates methods re-implemented on the CLIP architecture for fair comparison.

		Computation & N	Perfori	%)			
Method	# Params	Additional Storage	Peak GPU (GB)	GPU-Hours	Transfer	Avg.	Last
Storage-based Models							
ZSCL (Zheng et al., 2023)	149.6M	Data&Model (10.5GB)	43.1	47.24	68.1	75.4	83.6
MoE-Adapters (Yu et al., 2024)	59.8M	Routers (4.8GB)	15.5	3.42	68.9	76.7	85.0
DIKI (Tang et al., 2024)	1.8M	Task Stats (159MB)	10.2	4.40	68.7	76.3	85.1
InfLoRA [†] (Liang & Li, 2024)	7.8M	Grad. Proj. Mem. (9MB)	6.6	4.29	66.2	74.2	83.6
Storage-Free Models							
Continual-FT	149.6M	None	14.6	12.76	44.6	55.9	77.3
LoRA [†] (Hu et al., 2021)	15.7M	None	6.7	1.21	63.9	70.1	79.9
MiLoRA [†] (Wang et al., 2025)	15.7M	None	6.7	1.24	62.8	68.7	77.4
NuSA-CL (Ours)	1.5M	None	6.6	1.21	68.6	75.1	82.8

Table 2: Transfer, Avg., and Last (%) for Storage-free PEFT continual learning methods on the *5-shot* MTIL benchmark. † indicates methods re-implemented on the CLIP architecture for fair comparison.

	craf	Caltechil	J1 _RJ9	<i>B B</i>	Euro5 A	Flowers	Food	151	· "dP"	& \$\	5U1739 ⁵	
	Aire	Calter	CIFIX	Ø,	FINO	Elo.	\$0	ME	Oxfo.	Co	en,	AVG.
Zero-shot CLIP	24.3	88.4	68.2	44.6	54.9	71.0	88.5	59.4	89.0	64.7	65.2	65.3
Transfer												
LoRA [†] (Hu et al., 2021)	-	85.9	63.0	42.3	39.5	55.9	80.5	63.7	76.0	49.5	57.1	60.4
MiLoRA [†] (Wang et al., 2025)	-	84.8	59.6	43.2	37.7	50.6	78.3	61.4	80.2	42.4	55.6	59.4
InfLoRA [†] (Liang & Li, 2024)	-	87.2	65.9	44.5	52.1	64.3	85.3	63.4	83.8	58.9	62.5	66.8
NuSA-CL (Ours)	-	88.2	67.5	43.3	55.8	65.3	86.2	62.8	84.9	62.2	64.7	68.1
Avg.												
LoRA [†] (Hu et al., 2021)	15.9	90.6	68.1	54.1	69.1	74.3	81.8	73.6	81.9	58.3	62.0	66.8
MiLoRA [†] (Wang et al., 2025)	13.6	89.4	66.8	53.8	66.4	73.5	81.4	71.9	83.5	56.2	62.5	66.0
InfLoRA [†] (Liang & Li, 2024)	18.7	91.0	73.0	55.4	67.8	78.2	86.1	72.7	85.1	61.1	63.4	68.9
NuSA-CL (Ours)	28.9	90.5	73.2	56.1	71.9	76.9	87.2	72.9	86.8	63.5	65.3	70.3
Last												
LoRA [†] (Hu et al., 2021)	21.3	89.6	65.3	58.0	76.8	83.9	83.2	90.4	85.6	67.4	72.1	72.2
MiLoRA [†] (Wang et al., 2025)	17.5	88.8	66.0	56.3	72.3	82.4	82.1	87.3	87.7	68.2	72.0	71.0
InfLoRA [†] (Liang & Li, 2024)	21.9	91.3	73.3	58.9	77.8	90.0	87.9	88.6	89.8	71.3	72.6	74.8
NuSA-CL (Ours)	27.2	90.2	74.0	59.7	81.3	85.9	88.9	90.2	92.0	69.1	71.2	75.4

5.2 RESULTS

NuSA-CL Demonstrates a Superior Efficiency-Performance Tradeoff. Table 1 presents our main results on the full-shot MTIL benchmark. As shown, storage-based PEFT methods like MoE-Adapters (Yu et al., 2024) achieve the highest final accuracy. However, this comes at a significant cost: MoE-Adapters requires nearly 60M parameters and an expanding router library, while ZSCL (Zheng et al., 2023) incurs a massive computational overhead of 47.24 GPU-Hours. In contrast, NuSA-CL establishes a new state-of-the-art within the practical and challenging storage-free setting, significantly outperforming other competitors like LoRA and MiLoRA. The key finding is that NuSA-CL achieves performance highly competitive with the storage-based SOTA while being orders of magnitude more efficient. Specifically, compared to MoE-Adapters, NuSA-CL uses 40x fewer parameters (1.5M vs. 59.8M), zero additional storage, less than half the peak GPU memory, and is nearly 3x faster (1.21 vs. 3.42 GPU-Hours). This result highlights a vastly superior performance-to-cost tradeoff, positioning NuSA-CL as a powerful and scalable solution.

NuSA-CL Excels in Data-Efficient, Few-Shot Learning. To further probe the robustness of our approach, we conduct a focused analysis on the challenging 5-shot MTIL benchmark, with results in Table 2. As established in Table 1, Storage-free PEFT strategies represent the most practical paradigm for efficient continual learning. We therefore focus our analysis on this category to determine which

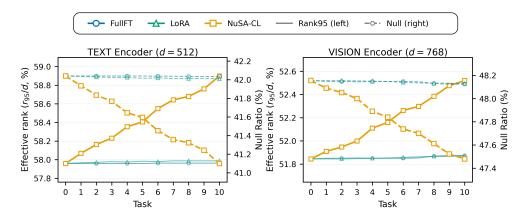


Figure 2: Our method actively accumulates knowledge by utilizing the model's intrinsic capacity, while other methods overwrite it. The plots track the evolution of the effective rank (solid lines, left axis) and the null ratio (dashed lines, right axis) for the text and vision encoders. NuSA-CL is the only method to show a steady increase in effective rank, demonstrating that it fills the underutilized null space to learn new tasks. This provides direct evidence for its ability to mitigate catastrophic forgetting.

low-rank adaptation strategy is most effective when data is scarce. This setting acts as a stress test, magnifying the fundamental differences between each approach. The results clearly demonstrate the superiority of our null-space adaptation strategy. NuSA-CL achieves the best performance across all summary metrics, decisively outperforming InfLoRA, the strongest competitor in this group. This indicates that our persistent null-space constraint is a fundamentally more robust and data-efficient strategy for mitigating forgetting than alternatives like subspace initialization (MiLoRA) or gradient projection (InfLoRA), validating the core mechanism of NuSA-CL.

Scalability in Long-Sequence Incremental Learning. Finally, to address the critical question of long-sequence scalability, we evaluate NuSA-CL on the Class-Incremental CIFAR100 benchmark (Table 4). The advantage of our method becomes increasingly pronounced as the task sequence lengthens. In the most challenging 50-step scenario, NuSA-CL achieves a final 'Last' accuracy of 71.85%, significantly outperforming the strongest baseline, ZSCL, by a large margin of over 4.4%. This result provides strong empirical evidence that our dynamic, task-wise re-computation of the null space is an effective and scalable strategy for lifelong learning, confirming the longevity of our approach even after 50 sequential tasks.

6 ANALYSIS

In this section, we analyze the effectiveness of NuSA-CL from multiple angles. We first visualize and explain how NuSA-CL's learning dynamics—knowledge accumulation versus overwriting—fundamentally differ from conventional methods. We then establish why adapting within the null space is a superior strategy for continual learning. Finally, we experimentally validate our core mechanisms and justify the choice of key hyperparameters.

6.1 Null Space Dynamics: Accumulation vs. Overwriting

Figure 2 illustrates a fundamental divergence in the learning dynamics of NuSA-CL compared to conventional fine-tuning approaches. The plots track the model's **effective rank** (solid lines), representing the capacity used to encode core knowledge, and the **null ratio** (dashed lines), the remaining underutilized capacity. The effective rank is defined as the minimum percentage of dimensions required to capture 95% of the weight matrix's total spectral energy (r_{95}/d) .

Conventional methods like LoRA and Full-FT exhibit a "lazy learning" behavior. As shown in Figure 2 and detailed in Appendix Table 10, their spectral properties remain almost static across all 11 tasks. For instance, the effective rank of LoRA's vision output projection layer barely changes,

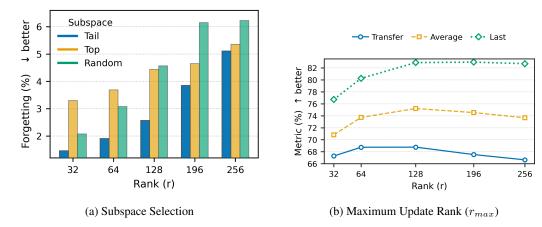


Figure 3: Ablations on subspace and rank. (Left) Subspace selection. Across ranks, Tail (null-like) consistently yields the lowest forgetting than Top and Random). (Right) Update rank. On the Tail subspace, overall CL performance peaks around $r_{\rm max} = 128$ (cf. Transfer/Avg./Last), indicating that a moderate rank balances retention (stability) and on-task adaptation (plasticity).

shifting trivially from an initial 447.42 to 447.58. This spectral inertia suggests that these methods do not exploit the model's underutilized capacity; instead, they primarily overwrite knowledge within the existing principal subspace, leaving the vast null space dormant.

In contrast, our method actively **accumulates knowledge** by progressively filling this underutilized space. For the same vision output layer, NuSA-CL's effective rank shows a clear and consistent increase. This trend, observed across all attention layers, provides direct quantitative evidence that NuSA-CL dynamically reshapes the parameter space to integrate new information. This additive learning process is the core mechanism behind NuSA-CL's ability to mitigate catastrophic forgetting and build a more informationally dense representation over time.

A natural question arises regarding the long-term viability of this approach: does the null space eventually become exhausted? Our analysis indicates that it does not. The "null space" is a low-energy spectral region, not a finite, empty container; it shrinks but is never fully depleted. The quantitative data in Appendix Table 10 confirms this. Even after learning 10 diverse and challenging tasks, the number of available null directions in the most saturated layer (vision output projection) is still 313.58. This is more than double our empirically chosen update rank ($r_{max}=128$), demonstrating that a substantial, low-interference subspace persists for future adaptation. This confirms the scalability and long-term robustness of our approach.

6.2 WHY THE NULL SPACE? SUBSPACE SELECTION STRATEGY

Tail outperforms Top and Random in mitigating forgetting. To identify an update subspace that mitigates catastrophic forgetting, we study how the choice of directions affects both knowledge retention and adaptation. We evaluate three initialization strategies for low-rank adaptation—Top (largest singular directions), Tail (smallest, null-like directions), and Random—using fixed rank r=128 and 1,000 training iterations per task on 11 MTIL datasets. We report Forgetting (average drop from post-task to final performance) in Fig. 3a and provide full per-rank numbers in Appx. Table 9. The Tail strategy, which exploits the null space, consistently yields the lowest forgetting across all ranks we tested, indicating that low-energy directions provide a safe region for updates with minimal interference to previously acquired knowledge. Quantitatively, Tail increases from 1.46% at r=32 to 5.11% at r=256, while remaining below both Top and Random at every rank (e.g., at r=128: Tail 2.57% vs. Top 4.44% and Random 4.57%). The mild rise with larger ranks suggests that a purer, lower-dimensional null subspace better preserves past information.

Balancing stability and plasticity for continual learning. Does this imply that "the smaller (purer) the null space, the better" for continual learning? Not necessarily. As shown in Fig. 3b, overall CL performance is maximized around an update rank of $r_{\rm max}$ =128, which balances stability (retention)

Condition

Train M & V_n

Train M, U_n , V_n

Train only M (Ours)

Table 3: Ablation studies validating NuSA-CL's design principles. (a) Core mechanisms such as the persistent constraint and multimodal adaptation are shown to be critical. (b) The method is practical, with negligible initialization overhead, and robust to hyperparameter choices.

(a) Core Mechanism Ablations

Transfer

68.58

66.37

62.60

Persistent Constraint Ab

lation		
Avg.	Last	
75.08	82.79	
73.11	82.04	
68.12	77.32	

	Modality Ablatio	n	
Modality	Transfer	Avg.	Last
Both (Ours)	68.58	75.08	82.79
Text-only	68.47	72.62	79.09
Vision-only	65.14	70.49	77.86

(b) Robustness and Practicality

Rob	Robustness to Cutoff Threshold (ρ)									
threshold	Transfer	Avg.	Last							
0.80	68.29	74.87	82.28							
0.90	68.82	75.07	82.74							
0.95	68.58	75.08	82.79							
0.99	68.49	74.89	82.70							
0.999	68.11	72.89	79.16							
	SVD Efficie	ency Analysis								
Method	Init. Time	Train Time (hr)	Avg.							

4.29

1.21

74.2

75.1

 \sim 81 min

<1 min

and plasticity (on-task adaptation). Consistent with Table 9, we observe a stability–plasticity tradeoff at *both* the update-rank dimension and the subspace choice: increasing rank improves **Target** performance but induces more forgetting, and *Top* attains marginally higher **Target** accuracy on the current task yet suffers substantially larger forgetting. Because continual learning is ultimately constrained by retention, these results motivate our design in NuSA-CL: operate in the null space (Tail) and cap the update rank at T_{max} to maintain stability while preserving competitive adaptation.

InfLoRA

Ours

6.3 VALIDATION OF NUSA-CL'S DESIGN PRINCIPLES

Core Mechanisms are Critical. As shown in Table 3a, our core design choices are essential for success. The *persistent constraint* is vital; unfreezing the null space bases (U_n, V_n) to make them trainable leads to a significant drop in performance, confirming that a strict, persistent constraint is necessary to prevent forgetting. Similarly, *multimodal adaptation* is superior, as jointly updating both text and vision encoders is key to maintaining cross-modal alignment.

Practicality and Robustness. A potential concern for our method is the overhead of SVD and sensitivity to hyperparameters. However, our analysis shows NuSA-CL is both practical and robust, shown in Table 3b. The SVD initialization is exceptionally lightweight. While our data-agnostic SVD is a one-time calculation per task with negligible overhead, competing methods like InfLoRA require heavy, data-dependent computations that scale poorly to design subspace using training dataset before learning. Furthermore, results show that performance is remarkably stable across a wide range of energy cutoff thresholds, demonstrating that NuSA-CL does not require sensitive hyperparameter tuning.

7 Conclusion

This paper tackles the challenge of adapting vision-language models to evolving tasks without catastrophic forgetting and without the unsustainable resource costs of methods whose storage or parameter counts grow linearly with the number of tasks, rendering them impractical for lifelong learning. We introduce **NuSA-CL**, a memory-free framework based on intrinsic adaptation. NuSA-CL identifies underutilized null space directions and constrains low-rank updates to this subspace, integrating new knowledge while preserving pre-trained capabilities. The learned update is then merged into the base model, maintaining a fixed parameter budget.

Across benchmarks, our method delivers a superior performance-efficiency trade-off: it outperforms other storage-free methods and rivals resource-intensive, storage-based approaches at a fraction of the cost. Strong results on long task sequences validate its scalability and effectiveness for lifelong learning, positioning NuSA-CL as a practical solution for deploying adaptable vision—language models in resource-constrained settings.

Limitations and future work. NuSA-CL remains robust on sequences of up to 50 tasks with ViT-B, but its capacity under extreme lifelong settings where the available null space directions may saturate warrants further study. We also note that the SVD step, while negligible with our reduced SVD on ViT-B, could become a bottleneck for substantially larger models. Future work includes (i) quantifying sensitivity to task order and semantic relatedness, as highly correlated sequences may concentrate usage of specific null-space dimensions; and (ii) developing lightweight, more reversible integration strategies, enabling selective forgetting without relying on persistent external memory.

REFERENCES

- Lukas Bossard, Matthieu Guillaumin, and Luc Van Gool. Food-101-mining discriminative components with random forests. In *Computer Vision–ECCV 2014: 13th European Conference, Zurich, Switzerland, September 6-12, 2014, Proceedings, Part VI 13*, pp. 446–461. Springer, 2014.
- Arslan Chaudhry, Marcus Rohrbach, Mohamed Elhoseiny, Thalaiyasingam Ajanthan, Puneet K. Dokania, Philip H. S. Torr, and Marc'Aurelio Ranzato. On tiny episodic memories in continual learning, 2019. URL https://arxiv.org/abs/1902.10486.
- Mircea Cimpoi, Subhransu Maji, Iasonas Kokkinos, Sammy Mohamed, and Andrea Vedaldi. Describing textures in the wild. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pp. 3606–3613, 2014.
- Li Deng. The mnist database of handwritten digit images for machine learning research [best of the web]. *IEEE signal processing magazine*, 29(6):141–142, 2012.
- Yuxuan Ding, Lingqiao Liu, Chunna Tian, Jingyuan Yang, and Haoxuan Ding. Don't stop learning: Towards continual learning for the clip model. *arXiv preprint arXiv:2207.09248*, 2022.
- Li Fei-Fei, Rob Fergus, and Pietro Perona. Learning generative visual models from few training examples: An incremental bayesian approach tested on 101 object categories. In 2004 conference on computer vision and pattern recognition workshop, pp. 178–178. IEEE, 2004.
- Patrick Helber, Benjamin Bischke, Andreas Dengel, and Damian Borth. Eurosat: A novel dataset and deep learning benchmark for land use and land cover classification. *IEEE Journal of Selected Topics in Applied Earth Observations and Remote Sensing*, 12(7):2217–2226, 2019.
- Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-Rank Adaptation of Large Language Models, October 2021.
- Moo Jin Kim, Karl Pertsch, Siddharth Karamcheti, Ted Xiao, Ashwin Balakrishna, Suraj Nair, Rafael Rafailov, Ethan Foster, Grace Lam, Pannag Sanketi, Quan Vuong, Thomas Kollar, Benjamin Burchfiel, Russ Tedrake, Dorsa Sadigh, Sergey Levine, Percy Liang, and Chelsea Finn. Openvla: An open-source vision-language-action model, 2024. URL https://arxiv.org/abs/2406.09246.
- Jonathan Krause, Michael Stark, Jia Deng, and Li Fei-Fei. 3d object representations for fine-grained categorization. In *Proceedings of the IEEE international conference on computer vision workshops*, pp. 554–561, 2013.
- Alex Krizhevsky, Geoffrey Hinton, et al. Learning multiple layers of features from tiny images. 2009.
- Zhizhong Li and Derek Hoiem. Learning without forgetting. *IEEE transactions on pattern analysis and machine intelligence*, 40(12):2935–2947, 2017.
- Yan-Shuo Liang and Wu-Jun Li. InfLoRA: Interference-Free Low-Rank Adaptation for Continual Learning, April 2024.
- Vijay Lingam, Atula Tejaswi Neerkaje, Aditya Vavre, Aneesh Shetty, Gautham Krishna Gudur, Joydeep Ghosh, Eunsol Choi, Alex Dimakis, Aleksandar Bojchevski, and Sujay Sanghavi. SVFT: Parameter-Efficient Fine-Tuning with Singular Vectors. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, November 2024.

- Haotian Liu, Chunyuan Li, Qingyang Wu, and Yong Jae Lee. Visual instruction tuning, 2023. URL https://arxiv.org/abs/2304.08485.
- Haotian Liu, Chunyuan Li, Yuheng Li, and Yong Jae Lee. Improved baselines with visual instruction tuning, 2024. URL https://arxiv.org/abs/2310.03744.
- Yue Lu, Shizhou Zhang, De Cheng, Yinghui Xing, Nannan Wang, Peng Wang, and Yanning Zhang. Visual Prompt Tuning in Null Space for Continual Learning, October 2024.
- Subhransu Maji, Esa Rahtu, Juho Kannala, Matthew Blaschko, and Andrea Vedaldi. Fine-grained visual classification of aircraft. *arXiv preprint arXiv:1306.5151*, 2013.
- Fanxu Meng, Zhaohui Wang, and Muhan Zhang. PiSSA: Principal Singular Values and Singular Vectors Adaptation of Large Language Models. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, November 2024.
- Maria-Elena Nilsback and Andrew Zisserman. Automated flower classification over a large number of classes. In 2008 Sixth Indian conference on computer vision, graphics & image processing, pp. 722–729. IEEE, 2008.
- Omkar M Parkhi, Andrea Vedaldi, Andrew Zisserman, and CV Jawahar. Cats and dogs. In 2012 IEEE conference on computer vision and pattern recognition, pp. 3498–3505. IEEE, 2012.
- Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, et al. Learning transferable visual models from natural language supervision. In *International conference on machine learning*, pp. 8748–8763. PMLR, 2021.
- Sylvestre-Alvise Rebuffi, Alexander Kolesnikov, Georg Sperl, and Christoph H Lampert. icarl: Incremental classifier and representation learning. In *Proceedings of the IEEE conference on Computer Vision and Pattern Recognition*, pp. 2001–2010, 2017.
- Gobinda Saha, Isha Garg, and Kaushik Roy. Gradient Projection Memory for Continual Learning, March 2021.
- Mustafa Shukor, Dana Aubakirova, Francesco Capuano, Pepijn Kooijmans, Steven Palma, Adil Zouitine, Michel Aractingi, Caroline Pascal, Martino Russi, Andres Marafioti, Simon Alibert, Matthieu Cord, Thomas Wolf, and Remi Cadene. Smolvla: A vision-language-action model for affordable and efficient robotics, 2025. URL https://arxiv.org/abs/2506.01844.
- Yuwen Tan, Qinhao Zhou, Xiang Xiang, Ke Wang, Yuchuan Wu, and Yongbin Li. Semantically-shifted incremental adapter-tuning is a continual vitransformer, 2024. URL https://arxiv.org/abs/2403.19979.
- Longxiang Tang, Zhuotao Tian, Kai Li, Chunming He, Hantao Zhou, Hengshuang Zhao, Xiu Li, and Jiaya Jia. Mind the Interference: Retaining Pre-trained Knowledge in Parameter Efficient Continual Learning of Vision-Language Models, July 2024.
- Pengwei Tang, Yong Liu, Dongjie Zhang, Xing Wu, and Debing Zhang. LoRA-Null: Low-Rank Adaptation via Null Space for Large Language Models, March 2025.
- Hanqing Wang, Yixia Li, Shuo Wang, Guanhua Chen, and Yun Chen. MiLoRA: Harnessing Minor Singular Components for Parameter-Efficient LLM Finetuning, March 2025.
- Shipeng Wang, Xiaorong Li, Jian Sun, and Zongben Xu. Training Networks in Null Space of Feature Covariance for Continual Learning, March 2021.
- Zifeng Wang, Zizhao Zhang, Sayna Ebrahimi, Ruoxi Sun, Han Zhang, Chen-Yu Lee, Xiaoqi Ren, Guolong Su, Vincent Perot, Jennifer Dy, et al. Dualprompt: Complementary prompting for rehearsal-free continual learning. In *European Conference on Computer Vision*, pp. 631–648. Springer, 2022a.
- Zifeng Wang, Zizhao Zhang, Chen-Yu Lee, Han Zhang, Ruoxi Sun, Xiaoqi Ren, Guolong Su, Vincent Perot, Jennifer Dy, and Tomas Pfister. Learning to prompt for continual learning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 139–149, 2022b.

- Xiwen Wei, Guihong Li, and Radu Marculescu. Online-LoRA: Task-free Online Continual Learning via Low Rank Adaptation, November 2024.
- Mitchell Wortsman, Gabriel Ilharco, Jong Wook Kim, Mike Li, Simon Kornblith, Rebecca Roelofs, Raphael Gontijo Lopes, Hannaneh Hajishirzi, Ali Farhadi, Hongseok Namkoong, et al. Robust fine-tuning of zero-shot models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 7959–7971, 2022.
- Yichen Wu, Hongming Piao, Long-Kai Huang, Renzhen Wang, Wanhua Li, Hanspeter Pfister, Deyu Meng, Kede Ma, and Ying Wei. S-LoRA: Scalable Low-Rank Adaptation for Class Incremental Learning, January 2025.
- Jianxiong Xiao, James Hays, Krista A Ehinger, Aude Oliva, and Antonio Torralba. Sun database: Large-scale scene recognition from abbey to zoo. In 2010 IEEE computer society conference on computer vision and pattern recognition, pp. 3485–3492. IEEE, 2010.
- Yibo Yang, Xiaojie Li, Zhongzhu Zhou, Shuaiwen Leon Song, Jianlong Wu, Liqiang Nie, and Bernard Ghanem. CorDA: Context-Oriented Decomposition Adaptation of Large Language Models for Task-Aware Parameter-Efficient Fine-tuning, March 2025.
- Jiazuo Yu, Yunzhi Zhuge, Lu Zhang, Ping Hu, Dong Wang, Huchuan Lu, and You He. Boosting Continual Learning of Vision-Language Models via Mixture-of-Experts Adapters. In *2024 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 23219–23230, Seattle, WA, USA, June 2024. IEEE. ISBN 979-8-3503-5300-6. doi: 10.1109/CVPR52733.2024.02191.
- Gengwei Zhang, Liyuan Wang, Guoliang Kang, Ling Chen, and Yunchao Wei. Slca: Slow learner with classifier alignment for continual learning on a pre-trained model, 2023. URL https://arxiv.org/abs/2303.05118.
- Zhen Zhao, Zhizhong Zhang, Xin Tan, Jun Liu, Yanyun Qu, Yuan Xie, and Lizhuang Ma. Rethinking Gradient Projection Continual Learning: Stability / Plasticity Feature Space Decoupling. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 3718–3727, 2023.
- Zangwei Zheng, Mingyuan Ma, Kai Wang, Ziheng Qin, Xiangyu Yue, and Yang You. Preventing Zero-Shot Transfer Degradation in Continual Learning of Vision-Language Models, August 2023.

A DATASET & IMPLEMENTATION DETAILS

A.1 BENCHMARKS AND METRICS.

Our primary evaluation is conducted on the *Multimodal Task Incremental Learning (MTIL) benchmark* (Zheng et al., 2023), which requires a model to sequentially train on 11 tasks while maintaining CLIP's zero-shot performance. The benchmark comprises Aircraft (Maji et al., 2013), Caltech101 (Fei-Fei et al., 2004), CIFAR100 (Krizhevsky et al., 2009), DTD (Cimpoi et al., 2014), EuroSAT (Helber et al., 2019), Flowers (Nilsback & Zisserman, 2008), Food (Bossard et al., 2014), MNIST (Deng, 2012), OxfordPet (Parkhi et al., 2012), StanfordCars (Krause et al., 2013), and SUN397 (Xiao et al., 2010). We report three metrics after completing the entire sequence:

- Transfer: Measures zero-shot transfer capability. After training on task t, we evaluate on the test sets of all future, unseen tasks $t + 1, \ldots, 11$ and average the results.
- Avg.: The mean test accuracy across all 11 datasets, recorded immediately after training on each task.
- Last: The mean test accuracy of the final model on each task's test set, capturing performance degradation (forgetting).

To test for long-sequence scalability, we also evaluate on the standard *Class-Incremental CIFAR100* benchmark (Krizhevsky et al., 2009), following the setup in ZSCL (Zheng et al., 2023).

A.2 IMPLEMENTATION DETAILS.

All experiments use the CLIP ViT-B/16 backbone with adapters applied to every query, key, value, and output projection. We use the AdamW optimizer (learning rate 3×10^{-4} , weight decay 10^{-2} , $\beta_1=0.9,\,\beta_2=0.999$), with a linear learning rate warmup for the first 5% of iterations followed by a cosine decay schedule. For NuSA-CL, we identify the null-space using a cumulative energy ratio cutoff (99% for 5-shot, 95% for full-shot) and cap the maximum rank at $r_{\rm max}=128$.

In the 5-shot MTIL setting (Yu et al., 2024), we sample five examples per class and train for 500 iterations with a batch size of 1. We scale each low-rank update by α/\sqrt{r} with $\alpha=1$, apply a dropout rate of 0.25 to the adapter branches, and use label smoothing of 0.2. In the full-shot setting, we use all training samples, train for 1000 iterations, and use a scaling factor of $\alpha=2$, keeping all other hyperparameters identical. All experiments were conducted on a single NVIDIA RTX 3090 GPU.

A.3 BASELINES.

We categorize baselines into three groups.

The first, *Full Fine-Tuning Models*, includes *Continual-FT*, which naively fine-tunes all parameters, and *ZSCL* (Zheng et al., 2023), which uses knowledge distillation.

The second group, *Storage-based PEFT Models*, require additional storage. This includes *MoE-Adapters* (Yu et al., 2024) and *DIKI* (Tang et al., 2024). It also includes *InfLoRA* (Liang & Li, 2024); while originally proposed for ViTs, we adapt its gradient projection memory mechanism to the multimodal CLIP architecture for a direct conceptual comparison.

The final group, *Storage-Free PEFT Models*, operates without additional storage. This includes the standard *LoRA* (Hu et al., 2021) method and *MiLoRA* (Wang et al., 2025), an LLM adaptation method we implemented on CLIP due to the relevance of its subspace initialization strategy. Our proposed method, *NuSA-CL*, also belongs to this category.

For a rigorous and fair comparison across all LoRA like methods (LoRA, MiLoRA, InfLoRA, and NuSA-CL), we implemented them within an identical framework: adapters were applied to both vision and text encoders with a consistent rank, and the updated weights were merged into the backbone after each task.

Table 4: Class-incremental CIFAR100 results: Last and Avg. accuracies (%) for 10/20/50-step splits. Bold = best, underline = second best.

	10 s	teps	20 s	teps	50 steps		
Methods	Last	Avg	Last	Avg	Last	Avg	
CLIP (Radford et al., 2021)	65.92	74.47	65.74	75.20	65.94	75.67	
Continual-FT	53.23	65.46	43.13	59.69	18.89	39.23	
LwF (Li & Hoiem, 2017)	48.04	65.86	40.56	60.64	32.90	47.69	
iCaRL (Ding et al., 2022)	70.97	79.35	64.55	73.32	59.07	71.28	
LwF-VR (Ding et al., 2022)	70.75	78.81	63.54	74.54	59.45	71.02	
ZSCL (Zheng et al., 2023)	73.65	82.15	69.58	80.39	67.36	79.92	
NuSA-CL (Ours)	74.51	<u>80.25</u>	73.84	81.03	71.85	80.19	

Table 5: Accuracy (%) of our method on the MTIL benchmark (5-shot, 500 iterations). Metrics for the Transfer, Last and Avg. are shown in the rightmost column.

	Aircraft	Caltech101	CIFAR 100	DTD	EuroSAT	Flowers	Food	MNIST	OxfordPet	StanfordCars	SUN397	Metric
Transfer		88.2	67.5	43.3	55.8	65.3	86.2	62.8	84.9	62.2	64.7	68.1
Aircraft Caltech101 CIFAR100 DTD EuroSAT Flowers Food MNIST OxfordPet StanfordCars SUN397	33.4 28.6 31.1 30.0 29.2 28.2 28.1 27.7 27.3 27.0 27.2	88.2 91.8 91.4 90.9 90.8 91.1 90.8 90.6 89.6 89.8 90.2	68.2 66.8 76.1 75.6 75.1 74.4 74.0 73.9 73.5 73.6 74.0	44.4 41.9 43.7 62.1 62.0 61.5 60.9 60.8 60.1 60.3 59.7	56.2 53.4 57.3 56.2 82.4 80.9 81.6 81.0 80.1 80.0 81.3	66.1 53.1 70.2 68.7 68.2 88.1 86.7 87.2 85.5 86.5 85.9	87.4 82.9 86.9 87.0 86.3 86.6 88.6 88.5 88.2 88.3	57.3 57.3 62.5 66.4 66.2 65.2 64.8 90.7 90.6 90.6	85.9 72.6 86.1 87.5 86.6 86.7 86.6 87.1 91.8 91.7 92.0	63.9 62.1 62.1 62.5 62.2 62.0 61.9 61.2 69.1	65.8 64.0 66.7 66.0 65.2 64.9 64.3 64.4 61.9 64.1 71.2	75,4
Average	28.9	90.5	73.2	56.1	71.9	76.9	87.2	72.9	86.8	63.5	65.3	70.3

B ADDITIONAL EVALUATION RESULTS

B.1 CLASS-INCREMENTAL LEARNING (CIL) RESULTS

We further evaluate on the standard CIFAR100 class-incremental splits (Krizhevsky et al., 2009), following ZSCL (Zheng et al., 2023). The 100 classes are grouped into 10, 20, or 50 tasks (10, 5, or 2 classes per task respectively). As the number of tasks increases, catastrophic forgetting becomes more severe. We use a learning rate of 3×10^{-3} , maximum rank $r_{\rm max}=256$, dropout of 0.05, batch size 128, and no label smoothing.

Table 4 shows that, although CLIP's zero-shot predictions already exceed naïve fine-tuning (FT) and LwF (Li & Hoiem, 2017), those baselines exhibit severe forgetting as tasks increase. Even ZSCL(Zheng et al., 2023)'s Last accuracy falls below 68% in the 50-step split. By contrast, our method yields the best Last performance across all splits (74.5%, 73.8%, and 71.9% for 10/20/50 steps) while maintaining very competitive Avg. scores. Remarkably, even in the prolonged 50-step regime, our method continues to recompute and leverage a fresh low-interference subspace, preserving prior knowledge effectively. This demonstrates that our method not only scales to the 11-task MTIL setting but also remains robust across very long class-incremental sequences.

B.2 MTIL COMPLETE RESULTS

Table 5 reports the results on all eleven tasks in the 5-shot MTIL regime (500 iterations per task). Our method maintains strong zero-shot retention (Transfer = 68.1%), while achieving an Average accuracy of 70.3% and a Last accuracy of 75.4%. The per-dataset breakdown confirms that our method uniformly preserves performance: no task suffers a dramatic collapse, and gains over baselines appear across both domain-shifted benchmarks (e.g., EuroSAT, Flowers) and in-domain benchmarks (e.g.,

Table 6: Accuracy (%) of our method on the MTIL benchmark (full dataset, 1 000 iterations). The rightmost column shows the overall Transfer, Last, and Average metrics.

	Aircraft	Caltech101	CIFAR100	DTD	EuroSAT	Flowers	Food	MNIST	OxfordPet	StanfordCars	SUN397	Metric
Transfer		88.3	66.8	44.0	55.5	67.9	85.8	66.7	84.8	60.7	65.2	68.6
Aircraft Caltech101 CIFAR100 DTD EuroSAT Flowers Food MNIST OxfordPet StanfordCars	49.7 42.9 41.0 41.6 39.9 38.8 38.6 34.3 33.2 33.2	88.3 96.7 96.1 96.2 95.6 95.7 95.5 95.9 95.9	67.5 66.0 82.2 81.6 81.0 80.6 80.7 79.8 79.6	44.3 42.3 45.4 74.2 73.6 72.1 73.5 72.4 72.3 71.8	53.6 53.2 59.8 55.6 97.0 96.9 96.9 96.6 96.6	69.9 64.5 69.2 68.2 67.9 96.4 95.9 95.9 95.2 94.6	88.3 85.2 86.2 85.6 85.0 84.8 91.1 90.6 90.6	58.0 55.5 73.7 71.2 69.7 69.9 68.7 98.9 98.9	87.7 83.9 85.1 84.6 84.3 82.9 85.4 84.9 94.8	63.2 60.7 61.6 61.5 60.6 60.1 60.3 59.6 59.0 79.2	65.5 63.8 66.4 65.8 65.1 64.7 65.7 65.7 64.4	02.0
SUN397 Average	35.0 38.9	95.2 95.2	79.4 78.0	71.5 64.9	96.2 81.7	94.0 82.9	90.6	98.8 78.4	95.0 87.6	78.1 64.0	76.9 66.3	82.8 75.1

Table 7: Transfer, Avg., and Last (%) for continual learning methods on the Order-2 sequence of the 5-shot MTIL benchmark. † indicates methods re-implemented on the CLIP architecture.

	Cats	Food	MIST	OxfordP	zi Flowers	SUMBOY	Aircraft	Caltechio	did	Euro5A	i CIFARIO	M.c.
Zero-shot	64.7	88.5	59.4	89.0	71.0	65.2	24.3	88.4	44.6	54.9	68.2	65.3
Transfer												
LoRA [†] (Hu et al., 2021)		87.6	63.0	86.6	63.5	63.2	19.8	87.4	43.8	44.0	61.3	62.0
MiLoRA [†] (Wang et al., 2025)		88.1	62.9	87.4	87.4	62.6	18.3	86.8	41.0	45.1	59.2	61.4
InfLoRA [†] (Liang & Li, 2024)		88.2	58.8	84.1	65.4	65.4	20.7	87.7	44.2	49.3	66.9	62.9
NuSA-CL (Ours)		87.6	60.0	86.3	65.8	63.8	21.9	88.3	43.6	53.8	68.3	63.9
Avg.												
LoRA [†] (Hu et al., 2021)	55.8	80.3	86.5	84.3	72.5	66.3	21.9	88.4	48.3	50.1	62.4	65.2
MiLoRA [†] (Wang et al., 2025)	51.0	76.7	83.8	81.7	71.7	64.3	19.2	87.7	44.1	50.7	60.1	62.8
InfLoRA [†] (Liang & Li, 2024)	65.3	85.5	85.3	85.6	80.4	67.1	25.3	89.3	48.3	54.4	67.6	68.6
NuSA-CL (Ours)	66.3	87.6	84.1	89.9	78.5	67.3	27.1	89.4	47.5	58.5	68.9	69.6
Last												
LoRA [†] (Hu et al., 2021)	46.7	76.7	89.3	82.1	71.3	67.9	23.7	90.0	59.1	72.4	73.3	68.4
MiLoRA [†] (Wang et al., 2025)	28.6	63.1	69.5	73.0	58.2	61.6	14.0	87.7	49.1	71.6	69.0	58.7
InfLoRA [†] (Liang & Li, 2024)	60.4	83.3	90.0	86.6	87.1	69.6	30.7	91.2	59.4	76.7	73.7	73.5
NuSA-CL (Ours)	64.6	86.7	89.4	91.7	84.7	70.1	32.7	91.8	58.0	79.2	75.1	74.9

CIFAR100, MNIST). Table 6 presents the analogous results in the full-shot regime (1000 iterations per task), showing the same pattern of robust Transfer, Avg., and Last scores.

B.3 MTIL ORDER-2 RESULTS

To assess sensitivity to task ordering, Table 7 and Table 8 report 5-shot and full-shot MTIL results, respectively, under a different task sequence (Order-2). In both settings, our method again achieves the highest Transfer, Average, and Last metrics, matching the original ordering (Order-1). Crucially, Last accuracy remains above 80% even in this permuted protocol, confirming that our method mitigates forgetting regardless of task order. This order-agnostic stability underscores the general applicability of our approach.

Table 8: Transfer, Avg., and Last (%) for continual learning methods on the Order-2 sequence of the the full dataset MTIL benchmark. † indicates methods re-implemented on the CLIP architecture.

	Cats	Food	MAIST	OxfordP'	e ^t Flor ^{der}	, shika	Aircraf	Caltechic	DID	Finos A	S CIFARIO	BAG.
Zero-shot	64.7	88.5	59.4	89.0	71.0	65.2	24.3	88.4	44.6	54.9	68.2	65.3
Transfer												
LoRA [†] (Hu et al., 2021)		88.2	59.7	86.9	65.9	65.1	20.7	88.3	44.8	58.6	62.0	62.0
MiLoRA [†] (Wang et al., 2025)		88.2	61.9	87.7	63.3	64.5	19.1	87.6	47.0	63.4	62.6	62.6
InfLoRA [†] (Liang & Li, 2024)		88.2	62.1	86.7	65.2	66.0	20.9	88.3	45.1	65.2	63.4	63.4
NuSA-CL (Ours)		88.1	57.7	87.0	66.2	64.8	21.9	89.2	49.4	66.9	63.4	63.4
Avg.												
LoRA [†] (Hu et al., 2021)	71.6	85.2	91.8	91.4	77.9	70.2	18.6	91.5	53.3	51.0	60.9	69.4
MiLoRA [†] (Wang et al., 2025)	67.0	85.3	92.3	90.9	76.0	69.6	20.2	90.8	51.8	55.9	65.3	69.6
InfLoRA [†] (Liang & Li, 2024)	80.0	89.2	92.4	92.2	82.8	71.9	30.8	91.5	53.5	55.4	67.0	73.3
NuSA-CL (Ours)	73.8	89.7	91.2	92.1	84.4	70.8	32.9	91.6	51.3	58.0	68.2	73.1
Last												
LoRA [†] (Hu et al., 2021)	52.5	78.1	96.8	91.9	72.5	71.6	3.9	97.0	73.4	84.7	84.3	73.3
MiLoRA [†] (Wang et al., 2025)	42.5	77.1	97.7	87.5	65.6	69.6	13.2	96.1	72.5	93.2	84.2	72.6
InfLoRA [†] (Liang & Li, 2024)	75.1	87.2	98.4	93.6	88.1	75.5	30.7	96.4	74.8	95.5	84.7	82.2
NuSA-CL (Ours)	64.6	88.6	98.4	93.4	93.9	75.8	44.8	95.8	72.5	96.4	80.9	82.9

Table 9: Per-rank results by subspace type. Lower is better for *Forgetting*; higher is better for *Target*.

Rank	Subspace	Forgetting (%)	Target (%)
	Tail	1.46	78.05
32	Top	3.30	79.66
	Random	2.08	78.70
	Tail	1.91	81.99
64	Top	3.69	82.96
	Random	3.08	82.77
	Tail	2.57	85.22
128	Top	4.44	85.50
	Random	4.57	85.49
	Tail	3.85	86.46
196	Top	4.65	86.76
	Random	6.15	86.67
	Tail	5.11	87.34
256	Top	5.36	87.46
	Random	6.23	87.44

C FURTHER ABLATION

This section provides detailed results for the analysis in Section 6. Table 9 presents the numerical data corresponding to the observations in Figure 3a. Table 10 lists the specific dimension numbers plotted in Figure 2.

Table 10: Effective rank r_{95} and corresponding null directions (Null@95 = $d-r_{95}$) for CLIP, LoRA and NuSA-CL across encoders and projection parameters.

Method	Encoder	Param	r_{95}	Null@95
		q	279.67	232.33
	TEXT	k	284.83	227.17
	IEAI	V	311.33	200.67
CL ID		0	311.17	200.83
CLIP		q	354.08	413.92
	MICION	k	358.42	409.58
	VISION	V	432.75	335.25
		o	447.42	320.58
		q	279.75	232.25
	TEXT	k	284.92	227.08
	TEXT	v	311.58	200.42
I DA (AC I		0	311.33	200.67
LoRA (After Learning)		q	354.17	413.83
	MICION	k	358.75	409.25
	VISION	V	432.92	335.08
		o	447.58	320.42
		q	282.58	229.42
	TEVT	k	288.25	223.75
	TEXT	v	316.67	195.33
NuCA CI (After I comine)		o	318.75	193.25
NuSA-CL (After Learning)		q	357.75	410.25
	MICION	k	362.42	405.58
	VISION	V	438.83	329.17